

The Unique Games Theorem

OpenAI

September 23, 2026

Abstract

We prove the Unique Games Conjecture. For every fixed $\varepsilon, \delta \in (0, 1/2)$, we give a deterministic polynomial-time reduction from 3SAT to Unique Games over a fixed finite alphabet, with completeness at least $1 - \varepsilon$ and soundness at most δ .

Contents

1	Introduction	2
1.1	The theorem	2
1.2	Proofs, codes, and the quantifiers	2
1.3	Grassmann expansion and the completeness obstacle	3
1.4	The construction and proof	4
2	Conventions	6
3	A latent alphabet gadget	7
3.1	The quadratic block	8
3.2	Generic spaces of characters	10
3.3	Concatenation and its noise	12
3.4	Linear lifts and loss of rank	14
3.5	Quotients and enlargement	16
4	The outer questions and the permutation test	17
4.1	Weighted parity equations and distinct positions	18
4.2	Tuples, local answer copies, and homogeneous coordinates	19
4.3	Exact table keys and folding	20
4.4	The latent matrix test and completeness	21
5	Matrix extraction and Fourier decoding	22
5.1	The inverse shortcode theorem	22
5.2	The advice experiment and the decoding statement	24
5.3	Comparing the two matrix perturbations	25
5.4	Obtaining many useful row fibers	26
5.5	Transferring the chosen slice	27
5.6	Decoding the row advice	30
6	Clean coordinates and the soundness gap	33
6.1	Classical projection-game repetition	33
6.2	Vanishing advice on clean coordinates	33
6.3	The matrix-game gap	35
7	Completion of the reduction	36
7.1	Removing weights and repeated edges	36
7.2	The final reduction	37
8	Approximation consequences	39
8.1	Constraint satisfaction and Vertex Cover	39
8.2	Ordering constraints	41
8.3	Cut, deletion, and clustering consequences	42
A	Finite formulations of the consequences	43
A.1	A simple strong-left form of Unique Games	43
A.2	The finite ordering transfer	46

1 Introduction

A permutation constraint prescribes exactly one label at one endpoint for each label at the other. The Unique Games Conjecture asserts that systems of such constraints can be hard to approximate even when almost all constraints can be satisfied. We resolve this conjecture positively, with an explicit unweighted reduction from 3SAT.

1.1 The theorem

A *Unique Games instance* has a finite vertex set, a finite alphabet K , and a nonempty list of oriented constraints $e = (u_e, v_e, \pi_e)$, where π_e is a permutation of K . A labeling a satisfies e when $a(v_e) = \pi_e(a(u_e))$. Its value is the fraction of satisfied constraints, and $\text{val}(G)$ is the maximum value of a labeling. The instance is *explicit and unweighted* when the entire vertex and constraint lists are written out, including the table of each permutation. A simple bipartite instance has two disjoint vertex sets and at most one constraint on each pair of opposite vertices.

Theorem 1.1 (Unique Games Theorem). *For every fixed $\varepsilon, \delta \in (0, 1/2)$, there are an integer $s \geq 1$ and a deterministic polynomial-time reduction from 3SAT to explicit unweighted Unique Games instances G_φ over $K = \mathbb{F}_2^s$ such that*

$$\begin{aligned} \varphi \text{ satisfiable} &\implies \text{val}(G_\varphi) \geq 1 - \varepsilon, \\ \varphi \text{ unsatisfiable} &\implies \text{val}(G_\varphi) \leq \delta. \end{aligned}$$

The graph is simple and bipartite, and every constraint is a translation $a(v_e) = a(u_e) + c_e$ of K . The alphabet, the degree of the running-time polynomial, and its constants depend only on ε, δ .

Thus the result resolves the Unique Games Conjecture positively in its usual edge-value formulation [23, Conjecture 2.5]. The two errors are prescribed independently, before the alphabet is chosen. Binary translation constraints are a known equivalent formulation of the conjecture [24, Corollary 13 and Section 11.1]; the reduction here produces that form directly.

Theorem 1.1 supplies the Unique Games hypothesis in earlier hardness reductions for cut, covering, constraint-satisfaction, ordering, deletion, and clustering problems [24, 27, 36, 17, 9, 10]. The reductions and the resulting approximation thresholds are due to those authors. Section 8 gives the precise problem formulations, strict thresholds, and reduction qualifications, together with proofs of the consequences.

1.2 Proofs, codes, and the quantifiers

The central difficulty is to retain completeness near one when every label permits only one reply. The earlier connection between proof verification and approximation hardness was developed by Feige, Goldwasser, Lovász, Safra, and Szegedy [15], by Arora and Safra through proof composition [3], and by Arora, Lund, Motwani, Sudan, and Szegedy through the constant-query PCP theorem [2]. Bellare, Goldreich, and Sudan introduced the long code and folding for verifier composition [7, Section 3.3]. The long code of an answer a is the table $f \mapsto f(a)$, indexed by Boolean functions on the answer set. In Håstad's parity-hardness proof, Fourier analysis recovers an outer-game strategy from arbitrary proof tables that make his verifier accept sufficiently often, without first recovering nearby codewords. His nearly satisfiable parity gap is our starting hardness input [18].

Khot introduced the Unique Games Conjecture in 2002 and gave an algorithm based on semidefinite rounding [22]. Several algorithmic guarantees explain why its order of quantifiers matters. For a fixed target error, Trevisan combines semidefinite rounding with graph decomposition

to obtain a polynomial-time guarantee when the input error shrinks at the reciprocal-logarithmic rate in the instance size [41]. Charikar, Makarychev, and Makarychev use randomized semidefinite rounding to obtain expected value $1 - O(\sqrt{\epsilon \log q})$ from a q -label instance of value at least $1 - \epsilon$ [8]; this guarantee loses accuracy as the alphabet grows, whereas the conjecture permits a large fixed alphabet after the errors are chosen. Kolla’s spectral enumeration gives a running-time bound in terms of the dimension of a suitable spectral subspace of the label-extended graph, whose vertices are variable–label pairs [30]. Arora, Barak, and Steurer use spectral decomposition and enumeration to obtain near-satisfying solutions for sufficiently small fixed completeness error, with a guarantee independent of the fixed alphabet size [1]; their running-time bound is subexponential, not polynomial, for fixed positive error. These are distinct limits on the known guarantees.

Integrality gaps concern the strength of a relaxation. Khot and Vishnoi constructed Unique Games with semidefinite value arbitrarily close to one and integral value arbitrarily close to zero, even with triangle inequalities [29]. Barak, Gopalan, Håstad, Meka, Raghavendra, and Steurer introduced the short code, based on Reed–Muller codes, as a shorter substitute for the long code in applications including alphabet reduction and quantitative integrality gaps [5]. These constructions constrain particular relaxations; a hardness reduction requires a separate argument. Bafna, Barak, Kothari, Schramm, and Steurer obtain polynomial-time constant-value guarantees for nearly satisfiable translation games whose graphs have suitable low-degree sum-of-squares hypercontractivity certificates, including noisy hypercube and noisy Reed–Muller short-code graphs [4]. The rank-one matrix shortcode used below is a different graph, and our argument does not assume these certificates.

1.3 Grassmann expansion and the completeness obstacle

The closest preceding hardness breakthrough concerns 2-to-1 and 2-to-2 games with imperfect completeness. A 2-to-1 constraint is a projection with exactly two preimages per right label. A 2-to-2 constraint partitions the labels on each side into pairs and matches these pairs bijectively, allowing the four combinations within a matched pair. Khot, Minzer, and Safra’s STOC 2017 work proposed a Grassmann-graph reduction whose nonstandard soundness guarantee was conditional on a structural hypothesis [26]. Dinur, Khot, Kindler, Minzer, and Safra’s STOC 2018 work proposed a reduction with ordinary edge-satisfaction soundness, conditional on an agreement hypothesis [12]; their expansion work formulated the relevant structural problem and established partial cases [11]. Barak, Kothari, and Steurer connected expansion to the needed agreement test through a binary matrix representation [6]. The expansion theorem of Khot, Minzer, and Safra completed this line in 2018 [25].

These results established hardness for almost satisfiable 2-to-1 and 2-to-2 games and, as a consequence, a Unique Games gap with completeness close to one half and arbitrarily small soundness. Obtaining Unique Games completeness arbitrarily close to one was a separate problem.¹

The degree-two matrix shortcode of Barak, Kothari, and Steurer [6] is the formulation used here: its vertices are binary matrices and its steps add rank-one matrices. Their two-reply rule permits either the old answer or its translate by the update’s left factor; requiring equality alone gives a one-half completeness guarantee. They discussed higher-degree tests as a possible route to improving this completeness [6, Section 1.6]. The inverse statement we use, derived from Grassmann expansion, appears in Theorem 5.1. Ellis, Kindler, and Lifshitz later gave a shorter expansion proof using hypercontractivity on spaces of linear maps [14]; our external input remains the Khot–Minzer–Safra theorem. Our proof keeps the established degree-two inverse theorem and changes the noise using a nonlinear map. The following comparison explains that choice.

¹A separate companion proves hardness for exact 2-to-1 games with perfect completeness and arbitrarily small fixed soundness, over fixed alphabets [34, Theorem 1.1]. It is not an input to this proof.

1.4 The construction and proof

The uniform rank-one test illustrates the completeness problem. Let M be a binary ℓ -by- m matrix and fix a nonzero answer $z \in \mathbb{F}_2^m$. The evaluation function $f_z(M) = Mz$ satisfies

$$f_z(M + al^\top) - f_z(M) = a(l^\top z).$$

For independent uniform $a \in \mathbb{F}_2^\ell$ and $l \in \mathbb{F}_2^m$, this evaluation keeps its value with probability $(1 + 2^{-\ell})/2$, close to one half. Our test changes the noise while retaining the ordinary rank-one test as the structural tool for soundness.

For any prescribed $p > 0$, Lemma 3.1 fixes a rank threshold and, for every sufficiently large ℓ , constructs a binary space $\mathcal{V}$ containing $K = \mathbb{F}_2^\ell$, a map $C : \mathcal{V} \rightarrow K$, and a noise law μ on $\mathcal{V}$ such that, for $a \sim \mu$,

$$C(x + h) = C(x) + h \quad (h \in K), \quad \Pr[C(x + a) \neq C(x)] \leq p$$

where x is uniform and independent of a . At the same time, every family of linear observations whose restrictions to K have sufficiently large rank detects a with probability at least $1/8$. The rank threshold depends only on p and is chosen before ℓ . The construction uses quadratic blocks over a finite field: their nonlinear outputs become increasingly stable under recursion, while a rank potential controls the loss of linear observations. Section 3 proves this separation and realizes the output translations as the subspace K .

For the same nonzero z , let $P \in \text{Hom}(\mathbb{F}_2^m, \mathcal{V})$ be uniform. The nonlinear evaluation $P \mapsto C(Pz)$ compares the values

$$C(Pz) \quad \text{and} \quad C(Pz + a(l^\top z))$$

under the update $P \mapsto P + al^\top$. Here Pz is uniform in $\mathcal{V}$, independently of the latent noise a , and $l^\top z$ is an independent fair bit. Its equality test therefore fails with probability at most $p/2$.

We start the reduction from binary parity equations on three distinct variables. Each equation has four satisfying triples. A table on an equation tuple $U = (e_1, \dots, e_k)$ is an affine map from the product of their solution planes to $\mathcal{V}$, represented by a matrix with $m = 1 + 2k$ columns: one intercept column and two slope columns per equation. Whenever one equation block affects a table only through one actual bit, its reduced key records that variable and the entire induced affine map. A projected table and its pullback receive the same key. Grouping constant K -translates of keys into a vertex then folds the answers: a label at the representative determines every translated answer, and a shared key gives the same restored answer on both domains. The resulting Unique Games test checks equality of the restored answers on a table and its rank-one perturbation using the latent noise. A near-satisfying global parity assignment gives the ideal evaluation above on every tuple it satisfies, supplying near completeness. The key sharing adapts Khot and Safra's virtual tables [28, Section 1.1], while the translation orbits adapt the folding principle of Bellare, Goldreich, and Sudan [7, Section 3.3].

For soundness, suppose a labeling passes this test with probability at least 0.99. A matrix frequency is a tuple $g = (g_i)_{i=0}^{m-1} \in (\mathcal{V}^*)^m$, one functional per column, with character $(-1)^{\sum_i g_i(Pe_i)}$, where the e_i are the standard basis vectors of $\mathbb{F}_2^m$. Averaging the phase increment of a rank-one update over the uniform selector l gives

$$\begin{aligned} \mathbb{E}_{a \sim \mu, l} (-1)^{\sum_i l_i g_i(a)} &= \Pr_{a \sim \mu} [g_i(a) = 0 \text{ for every } i], \\ \mathbb{E}_{a \text{ uniform in } K, l} (-1)^{\sum_i l_i g_i(a)} &= 2^{-\dim \text{span}\{g_i|_K\}}. \end{aligned}$$

The first multiplier is nonnegative and, by latent detection, is at most $7/8$ when the restriction rank is large. High acceptance therefore forces a fixed amount of low-rank Fourier mass; the second multiplier turns this into positive acceptance under the ordinary uniform rank-one perturbation in K . Using a linear complement $\mathcal{V} = K \oplus W$ fixed with the latent data, write the table as $P = (M, T)$, where M and T are its K - and W -valued affine components, represented by matrices that include their constant columns. The ordinary perturbation changes M and fixes T . Thus, for a positive fraction of pairs (U, T) , the restored answer as a function of M often keeps its value under the ordinary test. For the remaining soundness argument, take k to be a sufficiently large integer cube, with the latent and shortcode parameters fixed. On each such pair, the inverse shortcode theorem supplies agreement with a target $Mz + u$, for some $z \in \mathbb{F}_2^m$ and $u \in K$, on a slice defined by at most a fixed number r_s of row equations and at most r_s column equations.

For an independent uniform linear map $A : K \rightarrow \mathbb{F}_2^{r_s}$, the visible advice about (M, T) is (A, T, AM) . Together with U , it specifies T and the row fiber $\{M' : AM' = AM\}$ of K -valued affine maps on the answer space of U , including their constant terms. An erasure argument proves that a fixed positive fraction of such data under unrestricted table sampling on U admit additional column equations and a target $Mz + u$ certifying agreement on a nonempty slice. For each such datum (U, A, T, AM) , choose one certificate by a fixed rule depending only on that visible datum and the fixed labeling. Folding then makes the certificate determine a short nonempty set of valid answers on U .

The advice experiment compares the equation tuple U with a question O that independently replaces each equation by one uniformly chosen variable with probability β . Let π be the resulting projection of actual answers. The referee samples a uniform table (M_O, T_O) on O and an independent map A with the law above. On actual answer spaces, it sends the visible data $(A, T_O \circ \pi, AM_O \circ \pi)$ with U , and (A, T_O, AM_O) with O . With the sparse choice of β below, the full pulled-back table law, after averaging over projection choices, is close to the unrestricted table law used for extraction. On a positive fraction of joint samples the chosen U -certificate transfers to a nonempty slice on O and retains agreement.

The two strategies are defined from their respective visible data. When its advice admits a certificate, the U -strategy samples uniformly from its short answer set; on other inputs it returns any valid answer. The O -strategy uses a Fourier decoder on its own row fiber $\{M' : AM' = AM_O\}$. In the analysis, the transferred certificate supplies Fourier mass proving that this decoder reaches the projected U -answer set with positive probability; the decoder itself does not receive that certificate. The independent U -sample then gives full answer agreement with probability at least a fixed $\gamma > 0$, independently of k . The sparse question comparison and Fourier decoding adapt the method of Khot and Safra [28]; the correlated advice and use of vanishing coordinates follow Khot, Minzer, and Safra [26, Sections 3.3–3.4], with later use of that framework by Minzer and Zheng [31, Section 3].

On a NO parity instance, every pair of advice strategies has agreement tending to zero. The choice $\beta = k^{-2/3}$ makes $k\beta^2 \rightarrow 0$, as needed for the table-law comparison, while $k\beta \rightarrow \infty$ leaves many singleton coordinates. A singleton is clean when every observed row has zero coefficient on its retained bit. In the analysis, condition on A , the clean positions, the original advice intercept on O , and the sampled questions and advice slopes at the other positions. The clean question pairs then retain the independent equation–variable law. For each value of these fixed data, reconstructing each original local input from its own clean questions defines strategies for the repeated outer game. Parallel repetition bounds their success on the clean coordinates; averaging over the number of those coordinates bounds the original advice agreement by $\exp(-ck^{1/3})$. This contradicts the decoded lower bound for large k .

For repetition we use the explicit projection-game bound of Dinur and Steurer [13, Corollary 1.2]. Raz established exponential decay for general two-prover games [39], and Holenstein simplified

the proof [20]; those bounds depend on the answer alphabet. Rao removed that dependence for projection games [38, Theorem 4]. The Dinur–Steurer bound is applied first on the clean coordinates above and finally after rounding the matrix-game weights and subdividing every edge into four edges. Its alphabet independence permits the final repetition exponent to be fixed before the latent alphabet, completing the explicit unweighted reduction.

Figure 1 displays the reduction and the two incompatible bounds that establish soundness.




Figure 1: The solid arrows describe the reduction; the dashed arrows describe its soundness argument. The decoding lower bound follows from high test acceptance, while the clean-coordinate upper bound follows from the NO parity gap and applies to every advice strategy. The constants $c, \gamma > 0$ are fixed before the tuple length k grows.

Section 2 fixes the game and probability conventions. Section 3 constructs the gadget, and Section 4 defines the permutation test. Sections 5 and 6 give the two incompatible bounds on the advice experiment. Section 7 completes the parameter choices and the explicit Unique Games reduction. Section 8 then derives the approximation consequences from the completed theorem and the cited reductions. Appendix A gives the simple strong-left formulation and the finite ordering reduction used there.

2 Conventions

Weighted games. Intermediate instances may have loops, repeated constraint occurrences, and positive rational weights summing to one. Their value is the maximum total weight satisfied by a labeling. Occurrences retain separate IDs even if their constraints coincide. All rational numbers are encoded by binary numerators and denominators.

A bipartite *projection game* has a question for each vertex on either side, finite answer sets, and a map from left answers to right answers on each constraint. A strategy chooses one answer to each question. Answer sets that depend on the question can be identified with fixed alphabets of the same size; we use this for the four satisfying triples of a parity equation. Each prover uses only its own question. Private or shared randomness independent of the questions gives a mixture of deterministic strategies and cannot increase value.

The parallel power $G^{\otimes t}$ samples t independent constraint occurrences, sends the ordered question tuples, and accepts if every coordinate accepts. Answers may depend on the prover's entire tuple. Repeated question IDs retain separate coordinate positions and separate local answer copies. This is the ordinary classical product game used by the repetition theorem.

Binary linear algebra and Fourier analysis. Except where a larger field is specified, all vector spaces and maps are finite-dimensional over $\mathbb{F}_2$. Write $\text{Hom}(E, F)$ for the space of linear maps and $E^* = \text{Hom}(E, \mathbb{F}_2)$ for the dual. A subspace $S \leq E^*$ restricts to its image $S|_K \leq K^*$ on a subspace $K \leq E$; its restriction rank is $\dim(S|_K)$.

For a real function f on a binary space E , use normalized Fourier coefficients

$$\hat{f}(g) = \mathbb{E}_{x \in E} f(x)(-1)^{g(x)}, \quad f(x) = \sum_{g \in E^*} \hat{f}(g)(-1)^{g(x)}, \quad \sum_g |\hat{f}(g)|^2 = \mathbb{E}_x |f(x)|^2.$$

These identities follow because a nonzero linear functional has equally many zero and one values, so its character has mean zero. The characters are therefore an orthonormal basis. On a space of matrices, this is the entrywise pairing with its dual matrix space; we specify tensor notation when defining the Fourier decoder.

Probability and finite constructions. Uniform sampling from a vector space includes zero. Samples are independent unless a dependence is stated. For finite probability laws,

$$\text{TV}(P, Q) = \frac{1}{2} \sum_x |P(x) - Q(x)|.$$

It bounds the difference of expectations of any $[0, 1]$ -valued function. We condition only on events of positive probability.

Random experiments describe finite weighted constraint lists. The reduction enumerates those lists deterministically. Every auxiliary dimension and tuple length is fixed before the input formula varies; the final proof checks both this parameter order and the binary encoding cost. Fourier witnesses and decoders are used only to bound game value and need not be efficient.

3 A latent alphabet gadget

The construction in this section separates a nonlinear observable from families of linear observables under the same additive noise. All vector spaces and duals in this section are over $\mathbb{F}_2$, unless a different field is indicated. Restriction of a subspace of a dual means its image under the restriction map.

Lemma 3.1 (Latent alphabet gadget). *For every $p_* > 0$ there is an integer r_* such that, for every sufficiently large integer ℓ , the following objects exist: a finite-dimensional binary space $\mathcal{V}$ containing an identified subspace $K = \mathbb{F}_2^\ell$, a map $C : \mathcal{V} \rightarrow K$, and a rational probability distribution μ on $\mathcal{V}$. They satisfy*

$$C(x + k) = C(x) + k \quad (x \in \mathcal{V}, k \in K), \quad (3.1)$$

and, for independent uniform $x \in \mathcal{V}$ and $a \sim \mu$,

$$\Pr[C(x + a) \neq C(x)] \leq p_*. \quad (3.2)$$

Moreover, every linear subspace $G \leq \mathcal{V}^*$ with $\dim(G|_K) \geq r_*$ satisfies

$$\Pr_{a \sim \mu} [\text{there exists } g \in G \text{ with } g(a) \neq 0] \geq \frac{1}{8}. \quad (3.3)$$

In particular, r_* is fixed before ℓ is chosen. For rational p_* , the threshold and all the finite data can be found by a terminating deterministic algorithm.

We first construct a gadget with a fixed smaller output space B , in which shifts induce translations of B . After identifying shifts that induce the same translation, this gadget will detect every linear family whose restrictions fill B^* . We then combine copies indexed by injections $B \rightarrow K$: a sufficiently high-rank family on K restricts onto all of B^* on most copies. This second step allows every sufficiently large dimension of K while keeping the rank threshold fixed.

The fixed- B construction begins with a quadratic block. On an independent uniform input, a uniform nonzero perturbation in one of its update spaces leaves the nonlinear output unchanged with probability θ . A suitably generic rank- r family of linear observations loses at most one dimension on the same update space, and a uniform choice of that space causes a loss with probability at most $3r\theta$. A harmonic potential converts this into an expected loss of at most 3θ per level. Random changes of coordinates make nongeneric families rare as they descend, even when later linear maps depend on earlier choices. Concatenation therefore makes the nonlinear change probability small while leaving a constant probability of positive rank. At the final leaf, a surviving family detects an independent uniform symbol with constant probability.

3.1 The quadratic block

Fix an integer $r_0 \geq 1$, whose size will eventually depend only on p_* . Write $H_0 = 0$ and $H_r = \sum_{j=1}^r 1/j$. Let $F = \mathbb{F}_{2^d}$, $q = 2^d$, and regard $X = B = F^3$ as binary spaces, each of dimension $b = 3d$. The integer d will be chosen sufficiently large in terms of r_0 .

Use the trace pairing to identify both X^* and B^* with F^3 : the vector $z \in F^3$ represents the character $y \mapsto \text{Tr}_{F/\mathbb{F}_2}(z \cdot y)$, where the dot denotes the field dot product. We abbreviate the trace by Tr . The trace pairing is nondegenerate, and $\text{Tr}(t^2) = \text{Tr}(t)$. Every element of F has a unique square root.

Define

$$Q(x_1, x_2, x_3) = (x_2x_3, x_1x_3, x_1x_2), \quad D(x, a) = Q(x + a) + Q(x) + Q(a). \quad (3.4)$$

The block observable is

$$C_{\text{blk}} : X \oplus B \longrightarrow B, \quad C_{\text{blk}}(x, y) = y + Q(x).$$

It is equivariant under the vertical copy of B : $C_{\text{blk}}(x, y + k) = C_{\text{blk}}(x, y) + k$. For $z \in B^*$, the scalar observation $z \circ C_{\text{blk}}$ changes by $z(k)$ under this shift. We shall compare it with linear functionals whose restriction to the vertical B is also z .

For a one-dimensional F -subspace $A = Fv$ of F^3 , put

$$D_A = \{y \in B : v \cdot y = 0\}, \quad U_A = \{(a, Q(a) + w) : a \in A, w \in D_A\} \subseteq X \oplus B. \quad (3.5)$$

These definitions are independent of the chosen nonzero generator v . There are $q^2 + q + 1$ such lines; set

$$\theta = \frac{1}{q^2 + q + 1}. \quad (3.6)$$

Lemma 3.2. *For every line A , the space U_A is binary linear of dimension b . For $a \in A$ and $x \in X$, one has $D(x, a) \in D_A$; for $a, a' \in A$, one has $D(a, a') = 0$. The trace annihilator of D_A in $B^* = F^3$ is A , and the spaces U_A together span $X \oplus B$.*

Proof. The three coordinates of $D(x, a)$ are

$$x_2a_3 + a_2x_3, \quad x_1a_3 + a_1x_3, \quad x_1a_2 + a_1x_2.$$

When $a = tv$, every term in $v \cdot D(x, a)$ occurs twice, so their sum is zero. When both arguments are multiples of v , each coordinate of the polar expression is zero. Thus $Q|_A$ is binary additive, and the parametrization of U_A by $A \oplus D_A$ is a binary linear bijection. Its dimension is $d + 2d = b$.

Every $z \in A$ annihilates D_A under the trace pairing. The annihilator has binary dimension $3d - 2d = d$, so it is exactly A . Finally, for each coordinate line $A = Fe_i$, the space U_A contains $(Fe_i, 0)$, because $Q(te_i) = 0$, as well as $\{0\} \oplus D_A$. These three coordinate lines and their vertical hyperplanes span $X \oplus B$. $\square$

The nonlinear response. The form of U_A makes the quadratic term of an update cancel. For $(a, Q(a) + w) \in U_A$,

$$C_{\text{blk}}((x, y) + (a, Q(a) + w)) + C_{\text{blk}}(x, y) = w + D(x, a).$$

The right-hand side lies in D_A by Lemma 3.2. The following calculation gives the exact law that the recursion will use.

Lemma 3.3 (Response to one block update). *Fix a line A . For independent uniform $(x, y) \in X \oplus B$ and $\zeta \in U_A \setminus \{0\}$,*

$$\Pr[C_{\text{blk}}((x, y) + \zeta) \neq C_{\text{blk}}(x, y)] = 1 - \theta.$$

Proof. First take ζ uniform in all of U_A , still independently of the uniform aggregate. Its parametrization $(a, Q(a) + w)$ has independent uniform $a \in A$ and $w \in D_A$. For every fixed x, a , the output difference $w + D(x, a)$ is uniform in D_A , so it is nonzero with probability $1 - q^{-2}$. The zero perturbation has mass q^{-3} and never changes the output. Conditioning on a nonzero perturbation therefore gives

$$\frac{1 - q^{-2}}{1 - q^{-3}} = \frac{q(q + 1)}{q^2 + q + 1} = 1 - \theta.$$

$\square$

The lemma concerns an independent uniform aggregate and uniform nonzero update. When the update comes from a lower level of the recursion, we will verify this conditional law before applying the lemma.

Linear restrictions. Fix a binary subspace $S \leq B^*$. A linear functional on $X \oplus B$ whose restriction to the vertical B is $z \in S$ has a unique form $(\gamma_z, z) \in X^* \oplus B^*$. Thus a binary-linear choice of one such extension for each $z \in S$ is specified by a binary linear map $\gamma : S \rightarrow X^*$. We ask which of these chosen extensions vanish when restricted to U_A . The restriction loses rank exactly when its kernel contains a nonzero element.

Lemma 3.4 (Kernel of a block restriction). *Let $S \leq B^*$ and let $\gamma : S \rightarrow X^*$ be binary linear. For each line A , define*

$$\kappa_{A, \gamma} : S \longrightarrow U_A^*, \quad \kappa_{A, \gamma}(z) = (\gamma_z, z)|_{U_A}.$$

Under the trace identification, write $\gamma_z(x) = \text{Tr}(g(z) \cdot x)$ for the associated binary linear map $g : S \rightarrow F^3$. For every nonzero $z \in S$,

$$\kappa_{A,\gamma}(z) = 0 \iff A = Fz \text{ and } z \cdot g(z) = \sqrt{z_1 z_2 z_3}.$$

Proof. If z lies in the kernel, it annihilates every vertical vector $(0, w)$ with $w \in D_A$. Lemma 3.2 therefore gives $z \in A$, and for $z \neq 0$ this means $A = Fz$. On the remaining vectors $(tz, Q(tz))$, $t \in F$, the functional has value

$$\text{Tr}(t g(z) \cdot z + t^2 z_1 z_2 z_3). \quad (3.7)$$

The three summands of $z \cdot Q(tz)$ are equal, and three equals one in characteristic two. Also

$$\text{Tr}(t^2 c) = \text{Tr}(t \sqrt{c}).$$

Nondegeneracy of the trace pairing shows that Equation (3.7) vanishes for every t exactly when $z \cdot g(z) = \sqrt{z_1 z_2 z_3}$. Together with $z \in A$, this condition is sufficient as well: the vertical part is already annihilated. $\square$

3.2 Generic spaces of characters

The kernel criterion suggests two sufficient controls. We shall use a family in which no field line contains two nonzero characters and only a few characters satisfy the alignment equation for any one choice of extensions. The uniformity over that choice will later allow a lift to depend on earlier random changes of coordinates.

Definition 3.5. A binary subspace $S \leq F^3$ of dimension $r \geq 1$ is *generic* if no two distinct nonzero elements of S are proportional over F , and if, for every binary linear map $g : S \rightarrow F^3$, at most $3r$ nonzero $z \in S$ satisfy

$$z \cdot g(z) = \sqrt{z_1 z_2 z_3}. \quad (3.8)$$

Fix one map $\gamma : S \rightarrow X^*$ before choosing a uniform line A , and let $r_A = \dim \text{im } \kappa_{A,\gamma}$, where $r = \dim S > 0$. Lemma 3.4 now bounds the loss. If S is generic, each field line contains at most one nonzero element of S , so the kernel has dimension at most one. For the fixed map g associated with γ , at most $3r$ characters satisfy Equation (3.8); hence at most $3r$ lines can cause any rank loss. Consequently

$$r - r_A \leq 1, \quad \Pr_A[r_A < r] \leq 3r\theta, \quad \mathbb{E}_A[H_r - H_{r_A}] \leq 3\theta. \quad (3.9)$$

The last bound uses $H_r - H_{r-1} = 1/r$. For an arbitrary rank- r space, a loss still requires $A = Fz$ for some nonzero $z \in S$. Therefore

$$\Pr_A[r_A < r] \leq (2^r - 1)\theta \leq 2^r \theta, \quad H_r - H_{r_A} \leq H_r. \quad (3.10)$$

These estimates concern rank under restriction. Detection of the eventual noise will follow only after a positive-rank family reaches a leaf.

It remains to show that spaces with the stated genericity are abundant.

Lemma 3.6. *For each fixed $r \geq 1$, a uniformly random binary r -dimensional subspace of F^3 is generic with probability tending to one as $d \rightarrow \infty$.*

Proof. Choose independent uniform $x_{ij} \in F$, for $1 \leq i \leq 3$ and $1 \leq j \leq r$, and parametrize a binary linear map by

$$z_i(t) = \sum_{j=1}^r x_{ij} t_j, \quad t \in \mathbb{F}_2^r.$$

Each nonzero input has uniform image in F^3 , so the probability that this map is not injective is at most $(2^r - 1)q^{-3}$. Conditioned on injectivity, its image is a uniform binary r -subspace: every such subspace has the same number of ordered binary bases. Distinct nonzero binary inputs are linearly independent also over F , so their images are independent uniform elements of F^3 . A union bound therefore shows that the nonproportionality condition fails with probability $O_r(q^{-2})$, apart from the event of noninjectivity.

It remains to prove the alignment condition simultaneously for all g . Fix a set E of more than $3r$ distinct nonzero inputs in $\mathbb{F}_2^r$. On an injective parametrization, the values of a binary linear g on a basis can be arbitrary vectors in F^3 . Thus alignment on all of E implies that the vector

$$c = (\sqrt{z_1(t)z_2(t)z_3(t)})_{t \in E}$$

lies in the F -column span of the matrix with $3r$ columns

$$M_{t,(i,j)} = t_j z_i(t), \quad t \in E, \quad 1 \leq i \leq 3, \quad 1 \leq j \leq r.$$

Conversely, membership supplies the basis values of such a g . In particular, this column-membership condition already accounts for all possible maps g .

First regard the x_{ij} as independent indeterminates and set $L = \mathbb{F}_2(x_{ij})$. We claim that c is not in the column span of M even after extending scalars to an algebraic closure $\bar{L}$. The following r^3 elements are linearly independent over L :

$$\sqrt{x_{1j}x_{2k}x_{3l}}, \quad 1 \leq j, k, l \leq r.$$

Indeed, after clearing denominators, a dependence would be a polynomial identity in the square roots of the indeterminates. The displayed monomials have distinct parity vectors, whereas coefficients polynomial in the x_{ij} have even exponents in those square roots. Terms from different parity vectors cannot cancel.

The expansion of c in these independent elements has coefficient vectors

$$v_{jkl} = (t_j t_k t_l)_{t \in E}.$$

Suppose that c belonged to the column span over $\bar{L}$. Every row vector over L annihilating M would then annihilate c . The independence just proved implies that it would annihilate every v_{jkl} . By double annihilation in the finite-dimensional space L^E , every v_{jkl} would belong to the column span over L itself.

For $h \geq 1$, let $V_h \leq L^E$ be the evaluation span of all positive-degree monomials in the binary coordinates $t_1, \dots, t_r$ of degree at most h . Repeated indices in v_{jkl} give all monomials of degrees one and two as well as three, because $t_i^2 = t_i$ on the evaluation set. The preceding conclusion and the form of the columns therefore imply

$$V_3 \subseteq \text{col}_L(M) \subseteq V_2 \subseteq V_3.$$

These three spaces are equal. Multiplication by any t_i sends V_2 into V_3 , so preserves their common value. Since this space contains the degree-one monomials, it contains evaluations of every positive-degree monomial. These evaluations span all of L^E : for any nonzero binary point, its usual indicator

polynomial is a product of factors t_i and $1 + t_i$ with at least one factor t_i , and thus has no constant monomial. Its restriction to E gives the corresponding point indicator. Consequently the column span would have dimension $|E| > 3r$, a contradiction.

To pass from this algebraic statement to finite fields, square the entries of both M and c . The resulting matrices have polynomial entries, and the appended column has entries $z_1(t)z_2(t)z_3(t)$. Let $h = \text{rank}_L M \leq 3r$. Squaring preserves rank over $\overline{L}$, so the augmented squared matrix has a nonzero minor of size $h + 1$. This minor is a nonzero polynomial of degree at most $2h + 3 \leq 6r + 3$. All unaugmented minors of size $h + 1$ vanish identically, so the unaugmented rank cannot increase on specialization. It follows that specialized column membership forces this particular nonzero polynomial to vanish. For uniform independent $x_{ij} \in F$, this event has probability at most $(6r + 3)/q$ by the polynomial zero bound [40, Corollary 1]; see also the earlier sparse-polynomial formulation in [42, Section 3.1].

If the alignment condition fails, it fails on some set E of size $3r + 1$. There are finitely many such sets depending only on r ; if $2^r - 1 \leq 3r$ there are none and the condition is automatic. A union bound, followed by conditioning on injectivity, proves the lemma. $\square$

Genericity therefore controls every possible map g . We now fix a field where nongeneric spaces are rare at every rank up to r_0 ; no linear maps need to be chosen at this stage. Choose a positive ε_0 with

$$H_{r_0} 2^{r_0} \varepsilon_0 \leq 1, \quad \varepsilon_0 < 1. \quad (3.11)$$

By Lemma 3.6, we may fix d sufficiently large that $b = 3d \geq r_0$ and, for every $1 \leq r \leq r_0$, a uniform rank- r subspace of B^* is nongeneric with probability at most ε_0 . In particular, there exists a generic rank- r_0 subspace, which we fix once and for all.

3.3 Concatenation and its noise

We keep the output space B fixed while concatenating the quadratic blocks. One noisy branch will be chosen at each level. The next lemma computes how fast the nonlinear change probability decays; we will then track the rank of linear restrictions along the same branch.

We construct binary input spaces P_t , shift subspaces $R_t \leq P_t$, linear surjections $\lambda_t : R_t \rightarrow B$, and maps $C_t : P_t \rightarrow B$ such that

$$C_t(u + h) = C_t(u) + \lambda_t(h) \quad (h \in R_t). \quad (3.12)$$

A shift $h \in R_t$ thus represents the output translation $\lambda_t(h)$; different shifts may represent the same translation. This form of the shift action is preserved by aggregation. At the end, quotienting by the shifts of value zero will turn the output translations into an actual subspace of the input.

At height zero take $P_0 = R_0 = B$ and $C_0 = \lambda_0 = \text{id}_B$. Let $\mathcal{I}$ consist of all pairs (A, J) , where A is a line in F^3 and $J : B \rightarrow U_A$ is a binary linear isomorphism. Each line has the same number of such isomorphisms.

We include one child for every (A, J) in a single fixed parent map. Choosing a uniform noisy child will then sample a uniform line and a fresh uniform isomorphism. For $t \geq 1$, take $P_t = \prod_{(A, J) \in \mathcal{I}} P_{t-1}$ and set

$$(x, y) = \sum_{A, J} J C_{t-1}(u_{AJ}), \quad C_t(u) = C_{\text{blk}}(x, y) = y + Q(x). \quad (3.13)$$

A tuple of child shifts changes this aggregate by the sum of their embedded output translations. To make the parent change a pure output translation, we require the X -component of this sum to

vanish. Accordingly, define

$$R_t = \left\{ (h_{AJ}) \in \prod_{A,J} R_{t-1} : \sum_{A,J} J \lambda_{t-1}(h_{AJ}) \in \{0\} \oplus B \right\}, \quad (3.14)$$

and let $\lambda_t(h) = k$ when this aggregate is $(0, k)$. These are linear definitions, and Equation (3.12) follows immediately from Equation (3.13). Surjectivity follows inductively from Lemma 3.2: the child output translations can be chosen arbitrarily, and their aggregate map is onto $X \oplus B$. Equation (3.12) and surjectivity also show that C_t sends a uniform input to a uniform element of B .

Let ξ_t be the following additive noise in P_t . At height zero it is uniform in B . At a positive height, choose a uniform child (A, J) and embed an independent copy of ξ_{t-1} in that child, with zero in all other children. Equivalently, choose a uniform leaf in the product tree and add a uniform symbol of B at that leaf. Each path choice can be sampled as a uniform A followed by a uniform J , independently at successive levels. All these distributions are rational. Figure 2 illustrates one level of the recursion and the selected noisy branch.




Figure 2: One level of the latent recursion, shown under its additive noise. The input space is a product over all child indices (A, J) ; only one uniformly chosen child is perturbed. Each child output is embedded in $U_A \leq X \oplus B$, all embedded outputs are added, and the quadratic map is applied once at the parent. Here (x, y) is the initial aggregate, and Δ is the selected child's output difference, so the noisy aggregate is $(x', y') = (x, y) + J_*\Delta$. The omitted children obey the same rule. The highlighted perturbation may disappear at the parent.

Lemma 3.7. *For independent uniform $u \in P_t$ and noise ξ_t ,*

$$p_t := \Pr[C_t(u + \xi_t) \neq C_t(u)] = (1 - q^{-3})(1 - \theta)^t. \quad (3.15)$$

Proof. The identity at height zero is immediate. Consider one higher level and condition on the selected child (A, J) . Write Δ for its output difference. All child copies have the same function and the same inner noise law, so the law of Δ is independent of the selected index.

We need also the conditional law of the child's initial output. For any fixed child perturbation a , Equation (3.12) gives

$$C_{t-1}(u + a + h) + C_{t-1}(u + h) = C_{t-1}(u + a) + C_{t-1}(u) \quad (h \in R_{t-1}).$$

Thus the difference is invariant under these shifts, while the initial output is translated by every element of B . On uniform input its initial output remains uniform conditional on the difference,

even if the inner perturbation is also specified. All other initial child outputs remain independently uniform. Because the aggregate map is onto $X \oplus B$, the initial aggregate (x, y) is therefore uniform conditional on (A, J, Δ) . In particular it is independent of the aggregate perturbation $J\Delta$.

Conditional on $\Delta \neq 0$ and A , the isomorphism J remains uniform, because the law of Δ is independent of the selected child index. Thus $J\Delta$ is uniform in $U_A \setminus \{0\}$. The conditional uniformity just proved makes the initial aggregate independent of this update. Lemma 3.3 therefore gives parent change probability $1 - \theta$ conditional on a nonzero child difference. When $\Delta = 0$, the aggregate does not change. Consequently $p_t = p_{t-1}(1 - \theta)$, proving the identity. $\square$

3.4 Linear lifts and loss of rank

The nonlinear change probability now decays with the height. For the fixed- B gadget, our remaining target is detection whenever a family's restrictions fill B^* . We will prove this by tracking every binary-linear choice of extensions of these output characters through the shift action.

For $S \leq B^*$, a *lift* on P_t means a binary linear map $z \mapsto \psi_z$ from S to P_t^* satisfying

$$\psi_z|_{R_t} = z \circ \lambda_t. \quad (3.16)$$

A lift is *full* when its domain is B^* . We next describe every possible lift at a parent, without making any assumption about how it was chosen.

Lemma 3.8. *Let a lift with domain S be given on P_t , where $t \geq 1$. For each child (A, J) , restriction to that child's shift space factors through λ_{t-1} . The resulting characters on the child output space B have the form*

$$\phi_{AJ}(z) = J^*((\gamma_z, z)|_{U_A}) \in B^*, \quad (3.17)$$

where $\gamma : S \rightarrow X^*$ is a single binary linear map common to all the children.

Proof. A shift supported on one child and lying in $\ker \lambda_{t-1}$ is a parent shift with output translation zero. Each ψ_z vanishes on it, proving the first assertion.

Let $T = \prod_{A,J} B$ be the space of child output translations, with aggregate map $\mathcal{A} : T \rightarrow X \oplus B$. This map is onto. The sum of the child characters, minus z applied to the second component of $\mathcal{A}$, vanishes whenever the first component of $\mathcal{A}$ is zero: every such tuple of translations can be supplied by child shifts, giving a parent shift. The difference therefore factors uniquely through the first component of $\mathcal{A}$, which is onto X . Call this character γ_z . Uniqueness shows that γ_z depends linearly on z . Evaluating on a tuple supported at one child gives Equation (3.17). $\square$

After choosing a child, let $S' = \text{im } \phi_{AJ}$ and choose a linear section $\iota : S' \rightarrow S$ of ϕ_{AJ} . Restricting $\psi_{\iota(z')}$ to that child input gives a lift with domain S' on P_{t-1} . We may fix deterministic rules for all such sections. In particular they depend only on the already chosen path, and never on a future path index or on the final noise symbol. On descending the tree, every surviving functional is a restriction of a member of the original linear family.

The factorization now connects the recursion to the local restriction problem. For a fixed parent lift, its single map γ is common to every outgoing line A , and $\phi_{AJ} = J^* \kappa_{A,\gamma}$. Since J^* is an isomorphism, it does not change the kernel or rank. We may therefore apply Equations (3.9) and (3.10) when the outgoing line is chosen uniformly. The next proof shows that the incoming isomorphisms make genericity typical at each positive subsequent rank, even though the lift itself depends on the earlier path.

Proposition 3.9. *Set*

$$s = \left\lfloor \frac{H_{r_0}}{8\theta} \right\rfloor. \quad (3.18)$$

Every full lift on P_s detects ξ_s with probability at least $1/4$:

$$\Pr[\text{there exists } z \in B^* \text{ with } \psi_z(\xi_s) \neq 0] \geq \frac{1}{4}. \quad (3.19)$$

Proof. Restrict the full lift at the top to the fixed generic subspace of dimension r_0 . Follow the random leaf path, replacing the lift at each chosen child by the image lift just described. Let S_j be its domain after j steps, and let $r_j = \dim S_j$, so initially $r_j = r_0$ at $j = 0$. Continue with the zero space after extinction.

We spell out the conditioning that permits genericity to be used at every level. Let $\mathcal{F}_j$ be the information in all path choices before the outgoing choice (A_j, J_j) at step j . The current space, lift, and map γ_j are $\mathcal{F}_j$ -measurable. Conditional on $\mathcal{F}_j$ and A_j , the subspace

$$E_{j,A_j} = \text{im}[z \mapsto (\gamma_{j,z}, z)|_{U_{A_j}}] \leq U_{A_j}^*$$

is fixed. The next domain and its rank are

$$S_{j+1} = J_j^* E_{j,A_j}, \quad r_{j+1} = \dim E_{j,A_j}.$$

The rank is fixed before J_j is revealed. A uniform binary isomorphism $J_j : B \rightarrow U_{A_j}$ makes S_{j+1} a uniformly random subspace of B^* of that rank. Therefore, if B_j is the indicator that $r_j > 0$ and S_j is nongeneric, then

$$\mathbb{E}[B_{j+1} \mid \mathcal{F}_j, A_j] \leq \varepsilon_0. \quad (3.20)$$

The next lift may depend on J_j . This causes no restriction on the argument: genericity bounds every binary linear map g on the selected subspace, including the map supplied by that lift.

Let $\Delta_j = H_{r_j} - H_{r_{j+1}} \geq 0$. Conditional on $\mathcal{F}_j$, the outgoing A_j is independent and uniform, and the rank-loss estimates above apply to the fixed current lift. Consequently

$$\mathbb{E}[\Delta_j \mid \mathcal{F}_j] \leq 3\theta + H_{r_0} 2^{r_0} \theta B_j. \quad (3.21)$$

For $j = 0$ the space is generic, so its expected drop is at most 3θ . For $j \geq 1$, average Equation (3.21) over the preceding isomorphism using Equation (3.20). This gives the unconditional bound

$$\mathbb{E}\Delta_j \leq (3 + H_{r_0} 2^{r_0} \varepsilon_0) \theta \leq 4\theta.$$

No independence between an incoming orientation and its subsequent lift has been assumed.

Summing the expected drops and using Equation (3.18),

$$\mathbb{E}[H_{r_0} - H_{r_s}] \leq 4s\theta \leq \frac{H_{r_0}}{2}.$$

Extinction requires a total drop of H_{r_0} , so $\Pr[r_s = 0] \leq 1/2$. At a leaf, $P_0 = R_0 = B$ and its lift consists of the characters themselves. Conditional on any surviving rank $r_s > 0$, an independent uniform symbol in B is detected with probability $1 - 2^{-r_s} \geq 1/2$. Every surviving functional is a restriction of the original family along the selected leaf, so such a detection implies a detection by that family of the embedded noise ξ_s . Multiplying the two bounds proves Equation (3.19). $\square$

3.5 Quotients and enlargement

The fixed- B construction is complete once its shift action is realized by an actual subspace. Put

$$V_0 = P_s / \ker \lambda_s,$$

where the kernel is taken inside $R_s \leq P_s$, and let $q_0 : P_s \rightarrow V_0$ be the quotient map. Equation (3.12) makes C_s invariant under this kernel, so it descends to a map $\overline{C} : V_0 \rightarrow B$. The quotient $R_s / \ker \lambda_s$ is identified with B by λ_s , giving an inclusion $i_B : B \rightarrow V_0$ such that

$$q_0(h) = i_B(\lambda_s h) \quad (h \in R_s).$$

We henceforth identify this subspace with B . Its shift action satisfies

$$\overline{C}(x + b') = \overline{C}(x) + b' \quad (b' \in B).$$

Let ν be the pushforward of the noise law of ξ_s . Uniform input on P_s pushes to uniform input on V_0 , so the output change probability remains p_s .

Every linear subspace $G_0 \leq V_0^*$ whose restriction is all of B^* satisfies

$$\Pr_{a \sim \nu}[\text{there exists } g \in G_0 \text{ with } g(a) \neq 0] \geq \frac{1}{4}. \quad (3.22)$$

Indeed, choose a linear section $\sigma : B^* \rightarrow G_0$ of restriction and define $\psi_z = q_0^* \sigma(z) = \sigma(z) \circ q_0$. For $h \in R_s$,

$$\psi_z(h) = \sigma(z)(q_0 h) = z(\lambda_s h).$$

Thus this pullback is a full lift. Moreover, $\psi_z(\xi_s) = \sigma(z)(q_0 \xi_s)$, so detection by the lift is detection by the chosen section on the pushed-forward noise. Proposition 3.9 gives the bound, and detection by the section implies detection by G_0 .

Proof of Lemma 3.1. Choose r_0 sufficiently large that

$$\exp(1 - H_{r_0}/8) \leq p_*.$$

Then choose ε_0 as in Equation (3.11) and fix d , and hence $b = 3d$ and θ , using Lemma 3.6. Construct the height s in Equation (3.18) and the quotient $(V_0, B, \overline{C}, \nu)$ above. Lemma 3.7 gives

$$p_s \leq (1 - \theta)^s \leq \exp(-\theta s) \leq \exp(\theta - H_{r_0}/8) \leq p_*. \quad (3.23)$$

Set $r_* = b + 2$. All choices so far depend only on p_* .

We now enlarge the output space. The point of using random injections $B \rightarrow K$ is that a family of restriction rank at least $b + 2$ on K restricts onto all of B^* along most injections. A copy of the base gadget for each injection will therefore turn its full-restriction detection bound into the required high-rank bound.

Fix $\ell \geq b + 2$ and write $K = \mathbb{F}_2^\ell$. Let $\mathcal{L}$ be the finite set of all binary linear injections $L : B \rightarrow K$. Take a product $P = \prod_{L \in \mathcal{L}} V_0$ and define

$$\tilde{C}((x_L)_L) = \sum_{L \in \mathcal{L}} L \overline{C}(x_L).$$

The product of shift subspaces $T = \prod_L B \leq P$ acts on this output through the linear map

$$\Lambda : T \longrightarrow K, \quad \Lambda((b_L)_L) = \sum_L L b_L.$$

This map is onto, because every nonzero vector of K is in the image of some injection $B \rightarrow K$. Quotient by its kernel inside T :

$$\mathcal{V} = P / \ker \Lambda.$$

The map $\tilde{C}$ descends to $C : \mathcal{V} \rightarrow K$, and $T / \ker \Lambda$ identifies with the required shift subspace K . This proves Equation (3.1).

On P , choose a uniform $L \in \mathcal{L}$ and update only that copy by an independent sample from ν . Let μ be the resulting law after the quotient. On uniform product input the output sum changes exactly when the selected copy's output changes, because L is injective. Uniform input pushes to uniform input on the quotient, so the change probability is still p_s . Equation (3.23) proves Equation (3.2). Every noise distribution used here is obtained from uniform finite choices by pushforward, and therefore is rational.

It remains to establish detection for an arbitrary $G \leq \mathcal{V}^*$. Let $S = G|_K \leq K^*$, of dimension $r \geq r_*$. Pull G back to P and then restrict to its copy of V_0 indexed by L ; denote the resulting subspace of V_0^* by G_L . Its restriction to that copy's shift subspace B is exactly

$$G_L|_B = L^*S.$$

To see this, a shift $b' \in B$ supported on copy L represents the shift $Lb' \in K$ in the quotient, so a character restricting to $z \in S$ evaluates there as $z(Lb')$.

The map $L^* : S \rightarrow B^*$ fails to be onto if and only if there is a nonzero $b' \in B$ with $Lb' \in S^\perp$. For each such fixed b' , a uniform injection L makes Lb' uniform among the nonzero vectors of K . Hence

$$\Pr_L[Lb' \in S^\perp] = \frac{2^{\ell-r} - 1}{2^\ell - 1} \leq 2^{-r}.$$

A union bound gives

$$\Pr_L[L^*S \neq B^*] \leq (2^b - 1)2^{-r} \leq 2^{b-r} \leq \frac{1}{4}.$$

On each of the remaining copies, Equation (3.22) gives detection probability at least $1/4$ by G_L , and thus by G on the embedded noise. Averaging over the selected copy yields detection probability at least $\frac{3}{4} \cdot \frac{1}{4} = \frac{3}{16} \geq \frac{1}{8}$, proving Equation (3.3). The construction works for every $\ell \geq b + 2$, with the already fixed $r_* = b + 2$, as required.

For completeness, all choices are effective when p_* is rational. Choose an integer $m \geq 0$ with $2^{-m} \leq p_*$ and then choose r_0 with $H_{r_0} \geq 8(m + 1)$; since $e^{-m} \leq 2^{-m}$, this ensures the initial bound on r_0 . Choose rational ε_0 satisfying Equation (3.11). For successive integers d with $3d \geq r_0$, enumerate the binary subspaces of $\mathbb{F}_{2^d}^3$ and the linear maps in the definition of genericity. This tests genericity exactly and computes the fraction of nongeneric spaces at each rank $1, \dots, r_0$. Lemma 3.6 ensures that eventually all these fractions are at most ε_0 . Fix the first such field and a generic rank- r_0 space. The remaining products, isomorphisms, injections, quotients, function tables, and rational noise laws are obtained by finite enumeration and binary linear algebra. This gives the asserted terminating construction. $\square$

4 The outer questions and the permutation test

We now use the latent gadget to construct a permutation instance. Its vertices encode affine tables on tuples of parity equations. We identify two tables exactly when they induce the same affine map on the same retained coordinates, including the same intercept. This common representation will allow us to compare an equation tuple with a sparsely projected tuple in Section 5.

4.1 Weighted parity equations and distinct positions

A *weighted parity instance* consists of a finite variable set $\mathcal{X}$, a finite nonempty set $\mathcal{E}$ of equation occurrences, and positive rational weights ω_e summing to one. An occurrence is an ordered equation

$$x_{v_{e,1}} + x_{v_{e,2}} + x_{v_{e,3}} = b_e, \quad b_e \in \mathbb{F}_2.$$

Its value is the weighted fraction of equations satisfied, and $\text{OPT}(\mathcal{E})$ denotes the maximum value. Different occurrences have different IDs even if their ordered equations coincide.

Theorem 4.1 (Weighted parity gap). *For every fixed rational $\xi > 0$, there is a deterministic polynomial-time reduction from Boolean satisfiability to weighted parity instances, with binary-encoded rational weights, such that*

$$\begin{aligned} \text{YES: } & \text{OPT}(\mathcal{E}) \geq 1 - \xi, \\ \text{NO: } & \text{OPT}(\mathcal{E}) \leq \frac{1}{2} + \xi. \end{aligned}$$

The running-time polynomial may depend on ξ .

This is the gap consequence of Håstad's construction in [18, Theorem 5.4 and its proof]. More specifically, the proof uses a fixed dyadic verifier error $\delta > 0$, gives completeness $1 - \delta$ and soundness $(1 + \delta)/2$, and enumerates verifier choices as equations with their rational probabilities as weights. Choosing $\delta \leq \xi$ gives the displayed statement. The enumeration is a deterministic reduction; the verifier's random choices describe the resulting weighted instance.

Lemma 4.2 (Making positions distinct). *From a weighted parity instance $\mathcal{E}$ one can deterministically construct, in polynomial time, an instance $\mathcal{E}'$ in which the three variable IDs of every occurrence are distinct. Its weights are positive rationals of polynomial bit length, and*

$$\text{OPT}(\mathcal{E}') \geq \text{OPT}(\mathcal{E}), \quad \text{OPT}(\mathcal{E}') \leq \frac{1 + \text{OPT}(\mathcal{E})}{2} + \frac{3}{100}.$$

In particular, Theorem 4.1 with $\xi < 1/100$, followed by this preprocessing, gives completeness at least $1 - \xi$ and NO optimum at most $4/5$.

Proof. Replace each variable v by one hundred clones (v, c) , $c \in [100]$. For each original occurrence, choose its three clones uniformly subject to choosing distinct clones whenever positions have the same underlying variable. List all the resulting occurrences and give them the corresponding conditional weights. There are at most 100^3 choices per original occurrence. Giving all clones their original variable's label proves the completeness assertion.

For the upper bound fix arbitrary clone labels, and let

$$t_v = \frac{1}{100} \sum_{c=1}^{100} (-1)^{x_{(v,c)}}.$$

First choose clones independently, without conditioning on distinctness. The satisfaction probability of occurrence e is

$$\frac{1 + (-1)^{b_e} t_{v_{e,1}} t_{v_{e,2}} t_{v_{e,3}}}{2}.$$

Round each bias to a majority label, resolving zero biases arbitrarily. For every equation violated by this rounded assignment, the signed product in the numerator is nonpositive, so its satisfaction probability is at most $1/2$. The other equations have probability at most one. The weighted mean is

therefore at most $(1 + \text{OPT}(\mathcal{E}))/2$. The probability of a forbidden collision among the independently chosen clones is at most $3/100$, by a union bound over the three pairs of positions. Conditioning on its complement changes the expectation of any $[0, 1]$ -valued function by at most $3/100$. This proves the bound. All conditional probabilities have constant-size denominators apart from their original occurrence weights. $\square$

For the rest of the construction, $\mathcal{E}$ denotes the preprocessed instance, so every equation has three distinct variable IDs. Its *ordinary equation-versus-variable game* samples e according to ω and j uniformly from $[3]$. The first prover receives the occurrence ID e and answers with a satisfying triple for that equation; the second receives the variable ID $v_{e,j}$ and answers with a bit. Acceptance means equality at position j . Because the IDs within an equation are distinct, the question pair determines the projection constraint.

For any second-prover strategy, its bits constitute a global assignment. On an equation violated by that assignment, every valid first-prover answer disagrees in at least one position. Hence this ordinary projection game, denoted $G_{\mathcal{E}}$, satisfies

$$\text{val}(G_{\mathcal{E}}) \leq 1 - \frac{1 - \text{OPT}(\mathcal{E})}{3} \leq \frac{14}{15} \quad \text{if } \text{OPT}(\mathcal{E}) \leq \frac{4}{5}. \quad (4.1)$$

The answer alphabets have sizes four and two, respectively.

4.2 Tuples, local answer copies, and homogeneous coordinates

Fix a positive integer k . An equation question $U = (e_1, \dots, e_k)$ is sampled from $\omega^{\otimes k}$, with replacement. Its answer space is the product of the k affine solution planes. Blocks always use local copies: even if the same equation or variable ID appears in several blocks, their local answers are separate coordinates.

For a parameter $0 < \beta < 1$, independently in every block either retain the equation, with probability $1 - \beta$, or retain just the variable ID at a uniformly chosen position, with probability β . The resulting question is denoted O , and π is the coordinate projection from answers on U to answers on O . The question O records an equation ID or a variable ID at every ordered block position. A singleton block accepts either bit.

Represent an equation answer by its first two bits (x_i, y_i) ; the third is $b_{e_i} + x_i + y_i$. A singleton is represented by its actual bit. Add a common homogeneous coordinate t_0 . Thus

$$L_U = \mathbb{F}_2^{1+2k}, \quad m = 1 + 2k,$$

and L_O has one homogeneous coordinate, two coordinates for each retained equation, and one for each singleton. In either space, vectors whose first coordinate is one are exactly valid answers. Write $e_0^\top$ for the functional reading that coordinate.

The affine projection π extends to a linear surjection $\bar{\pi} : L_U \rightarrow L_O$ preserving t_0 . On retained equation blocks it is the identity, and on a block projected to a singleton its coordinate is

$$\begin{cases} x_i, & \text{position 1,} \\ y_i, & \text{position 2,} \\ b_{e_i} t_0 + x_i + y_i, & \text{position 3.} \end{cases} \quad (4.2)$$

We call every map obtained by these blockwise choices a *permitted projection*; this class of maps is independent of β . Surjectivity follows by fixing t_0 and choosing coordinates independently in each block. We use the same conventions for questions with some blocks omitted.

4.3 Exact table keys and folding

The identification of tables across projected questions follows the virtual-table approach of Khot–Safra [28, Section 1.1]. We first express each table on the coordinates through which it factors. This will give the same representation to a projected table and its pullback.

Fix latent data $(\mathcal{V}, K, C, \mu)$ from Lemma 3.1, where $K = \mathbb{F}_2^\ell$, and fix a linear complement $\mathcal{V} = K \oplus W$. A *table* on a question Q is an affine map from its answer space to $\mathcal{V}$, or equivalently its unique linear extension $P \in \text{Hom}(L_Q, \mathcal{V})$. Here Q may be an equation question or one of its permitted projections, and the same definition applies to a question with omitted blocks. Write c for the coefficient of the homogeneous coordinate t_0 in P .

In an equation block, inspect the rank of its slope map $\mathbb{F}_2^2 \rightarrow \mathcal{V}$, whose contribution we write as $a_i x_i + d_i y_i$. If its rank is two, retain the equation occurrence and this contribution. If its rank is zero, omit the block. In the remaining case the rank is one, so the coefficients have exactly one of the forms

$$(a_i, d_i) = (v, 0), \quad (0, v), \quad (v, v), \quad v \neq 0.$$

These forms retain only the variable at position one, two, or three, respectively, with contribution $v z_i$ in its actual bit z_i . For the third form,

$$v(x_i + y_i) = v z_i + b_{e_i} v t_0, \quad z_i = b_{e_i} t_0 + x_i + y_i,$$

so this replacement also adds $b_{e_i} v$ to the homogeneous intercept. In a singleton block, retain its variable ID and its original contribution exactly when its slope is nonzero; otherwise omit the block.

The retained blocks form the *support* of (Q, P) . They keep their indices in $[k]$, their equation/variable tags, and their IDs. A retained equation keeps its ordered first two coordinates, and a retained singleton uses its actual-bit coordinate. Thus repeated IDs in different blocks still have separate local coordinates. The support may be empty; its homogeneous coordinate remains in every case.

Let J_3 be the set of equation blocks reduced to their third bit, and write v_i for the nonzero slope in such a block. The *reduced table* $\tilde{P}$ on the support has the retained block contributions specified above and homogeneous intercept

$$\tilde{c} = c + \sum_{i \in J_3} b_{e_i} v_i. \quad (4.3)$$

To check the complete table, let L_{supp} be the homogeneous answer space of the support and let $\rho_{Q,P} : L_Q \rightarrow L_{\text{supp}}$ retain t_0 and the specified block coordinates. It is the identity on retained equations and on retained singleton blocks already present in Q , uses the corresponding actual-bit coordinate from Equation (4.2) on a reduced equation, and discards omitted blocks. Choosing the coordinates independently in each block shows that $\rho_{Q,P}$ is surjective. The block identities and Equation (4.3) give

$$P = \tilde{P} \rho_{Q,P}.$$

Surjectivity makes $\tilde{P}$ the unique linear map with this factorization. We define the *key* of (Q, P) to be its support together with this entire reduced table. Thus two tables have the same key exactly when their retained coordinates agree, with the same positions, tags, and IDs, and their induced affine maps agree on those coordinates, including their intercepts.

We next check compatibility with every permitted projection $\bar{\pi} : L_U \rightarrow L_O$. Let $P' \in \text{Hom}(L_O, \mathcal{V})$. An equation retained in O has the same contribution in P' and its pullback $P' \bar{\pi}$, so its reduction is identical. A singleton of slope zero disappears on both sides. A singleton of nonzero slope v pulls back to the corresponding one of the three rank-one patterns above, which reduces to the same variable ID at the same block position. For a third-position singleton, the pullback adds $b_{e_i} v$ to the

intercept, and its reduction adds $b_{e_i}v$ once more. The two terms cancel over $\mathbb{F}_2$. All other intercept corrections come from retained equation blocks and are identical on the two sides. The supports and complete reduced tables therefore agree: $(U, P'\bar{\pi})$ and (O, P') have the same key.

Define the key universe to consist of the keys of all tables on equation questions $U \in \mathcal{E}^k$. The compatibility just proved shows that the key of every table on a permitted projected question already belongs to this universe, by pullback from any compatible U .

We now apply the folding principle of Bellare, Goldreich, and Sudan [7, Section 3.3]: values on one representative determine the values on its translates. Adding a constant $h \in K$ to a key means adding h to its reduced intercept, leaving its support and slopes fixed. This is realized by replacing a presenting table P with $P + he_0^\top$, so it preserves the key universe. The action is free because the key records the intercept. A vertex of our permutation instance is an orbit of this action, represented by its unique key whose reduced intercept has zero K component.

For a table P on an equation question or a permitted projected question Q , let $c(Q, P) \in K$ be the K component of its key's reduced intercept. Let $v(Q, P)$ denote the orbit vertex represented by the key obtained after removing this component. A labeling λ of the vertices defines the restored table answer

$$F_Q(P) = \lambda(v(Q, P)) + c(Q, P).$$

Adding a constant changes only the removed component, so

$$F_Q(P + he_0^\top) = F_Q(P) + h \quad (h \in K). \quad (4.4)$$

Tables with the same key have the same representative and removed component. The pullback compatibility therefore gives, for every permitted $U, O, \bar{\pi}$ and every table P' on O ,

$$F_U(P'\bar{\pi}) = F_O(P'). \quad (4.5)$$

Projected questions thus use answers determined by the same vertex labeling, without an additional consistency test.

4.4 The latent matrix test and completeness

The permutation instance, denoted $\mathcal{G}$, samples independently

$$U \sim \omega^{\otimes k}, \quad P \text{ uniform in } \text{Hom}(L_U, \mathcal{V}), \quad a \sim \mu, \quad l \text{ uniform in } \mathbb{F}_2^m,$$

and checks

$$F_U(P) = F_U(P + al^\top). \quad (4.6)$$

Here $al^\top$ denotes the map $z \mapsto a(l^\top z)$. Each sample specifies an edge between the two key-orbit vertices. If their restored offsets are c_0, c_1 , its constraint is the permutation

$$\lambda(v_1) = \lambda(v_0) + c_0 + c_1$$

of the fixed alphabet K . Edge occurrences have their sampling probabilities as weights. At this stage loops and repeated edges are allowed. The sparse projection experiment involving O is used to analyze these same vertex labels; it is not an additional test in $\mathcal{G}$.

Lemma 4.3 (Completeness of the outer matrix test). *If $\mathcal{E}$ has an assignment of value at least $1 - \xi$, and the latent map satisfies*

$$\Pr_{x \text{ uniform in } \mathcal{V}, a \sim \mu} [C(x + a) \neq C(x)] \leq p_*,$$

then

$$\text{val}(\mathcal{G}) \geq 1 - k\xi - \frac{p_*}{2}.$$

Proof. Fix the near-satisfying global bit assignment. Give every reduced support a valid answer as follows. At a singleton use its global bit. At a retained equation use its first two global bits and set the third to the unique value satisfying the equation. These choices depend only on the support, including its occurrence IDs, and therefore define an answer for every exact key independently of its origin.

At an orbit representative, evaluate its reduced table on this answer and apply C . This defines a vertex labeling. The equivariance of C under K implies that its restored answer is evaluation of the actual reduced table followed by C , in agreement with Equation (4.4).

With probability at least $1 - k\xi$, every equation in U is satisfied by the global assignment. On such a tuple, the support answers used for *every* table are projections of the same full answer $z_U \in L_U$, whose first coordinate is one. Consequently the two sides of Equation (4.6) are

$$C(Pz_U) \quad \text{and} \quad C(Pz_U + a(l^\top z_U)).$$

Since $z_U \neq 0$, the vector Pz_U is uniform in $\mathcal{V}$, independently of a, l , and $l^\top z_U$ is an independent fair bit. The test therefore fails with probability at most $p_*/2$ on these tuples. Adding the probability of a tuple containing an unsatisfied equation proves the bound. $\square$

5 Matrix extraction and Fourier decoding

We convert high acceptance of the matrix test into agreement between two strategies for the sparsely projected questions U, O . Each strategy uses only its own question and a bounded number of rows of matrix advice. The resulting agreement probability is positive independently of the tuple length k and the parity instance. The table functions enter only through their folding and sharing identities and the acceptance bound, so the proof applies to any family with these properties; the restored functions of every vertex labeling are examples. We state the precise assumptions below, after fixing the parameters and defining the advice experiment.

5.1 The inverse shortcode theorem

The structural input describes functions that often keep their value under a uniform rank-one perturbation. Such a function agrees with an affine evaluation on a set obtained by fixing a bounded number of linear combinations of rows and columns.

Write $\text{Mat}_{\ell, m} = \text{Hom}(\mathbb{F}_2^m, \mathbb{F}_2^\ell)$. A *row and column slice* is a nonempty set specified by equations

$$\mathcal{S} = \{M : d_i^\top M = s_i^\top \ (1 \leq i \leq n_r), \quad Mq_j = t_j \ (1 \leq j \leq n_c)\}, \quad (5.1)$$

where $d_i \in \mathbb{F}_2^\ell$, $s_i \in \mathbb{F}_2^m$, $q_j \in \mathbb{F}_2^m$, and $t_j \in \mathbb{F}_2^\ell$. The counts n_r, n_c need not be the ranks of their respective systems.

Theorem 5.1 (Equality soundness for degree-two shortcode). *For every $\eta \in (0, 1)$ there are $\alpha \in (0, 1]$, an integer $r_s \geq 1$, and an integer ℓ_0 , depending only on η , with the following property. For every $\ell \geq \ell_0$ there is $m_0 = m_0(\eta, \ell)$ such that, for every $m \geq m_0$ and every function $f : \text{Mat}_{\ell, m} \rightarrow \mathbb{F}_2^\ell$, if*

$$\Pr_{M, a, l} [f(M) = f(M + al^\top)] \geq \eta, \quad (5.2)$$

then some slice $\mathcal{S}$ as in Equation (5.1), with $n_r, n_c \leq r_s$, and some $z \in \mathbb{F}_2^m$, $u \in \mathbb{F}_2^\ell$ satisfy

$$\Pr_{M|\mathcal{S}} [f(M) = Mz + u] \geq \alpha. \quad (5.3)$$

Here M , $a \in \mathbb{F}_2^\ell$, and $l \in \mathbb{F}_2^m$ are independent and uniform, and $M \mid \mathcal{S}$ is uniform on the slice. No folding condition is imposed on f .

Proof. We derive this consequence from the Grassmann expansion theorem of Khot–Minzer–Safra [25, Theorem 1.12 in the 2018 revision], using the matrix chart of Barak–Kothari–Steurer [6, Definition 3.2 and Lemmas 3.3–3.7].

The Grassmann graph on an ambient binary space E has the ℓ -dimensional subspaces as vertices; two vertices are adjacent when their intersection has dimension $\ell - 1$. The *retention* of a nonempty vertex set is the probability that a uniform vertex in the set has a uniform neighbor still in the set. The cited theorem says that a set of retention at least a fixed $\zeta > 0$ has positive relative density in some interval

$$\{L : A_0 \subseteq L \subseteq B_0\}, \quad \dim A_0 + \operatorname{codim}_E B_0 \leq r.$$

The density bound and r depend only on ζ . The theorem holds for all sufficiently large ℓ , and then all sufficiently large $\dim E$.

Equality acceptance is the average of the retention probabilities of the nonempty fibers $f^{-1}(y)$ under the matrix step, weighted by their uniform masses. Some fiber S therefore has retention at least η . The probability that either perturbation factor is zero is

$$p_0 = 2^{-\ell} + 2^{-m} - 2^{-\ell-m}.$$

Choose the dimension thresholds so that $p_0 \leq \eta/2$. Conditional on both factors being nonzero, the retention of S is at least $(\eta - p_0)/(1 - p_0) \geq \eta/2$.

For each matrix put

$$L_M = \{(x, M^\top x) : x \in \mathbb{F}_2^\ell\} \subseteq E := \mathbb{F}_2^\ell \oplus \mathbb{F}_2^m.$$

The map $M \mapsto L_M$ is a bijection onto the chart $\mathcal{C}$ of subspaces whose projection onto the first summand is invertible. Since

$$\dim(L_M \cap L_N) = \ell - \operatorname{rank}(M - N),$$

the neighbors within the chart correspond exactly to nonzero rank-one matrix differences. Each chart vertex has $(2^\ell - 1)(2^m - 1)$ such neighbors, whereas its full Grassmann degree is $(2^\ell - 1)(2^{m+1} - 2)$: choose a hyperplane in the vertex and then a one-dimensional extension in the quotient by that hyperplane, other than the vertex itself. Thus exactly half of its neighbors lie in the chart. Over $\mathbb{F}_2$, every nonzero rank-one matrix has a unique factorization $al^\top$ with nonzero factors. The lifted set $\tilde{S} = \{L_M : M \in S\}$ consequently has Grassmann retention at least $\eta/4$.

Apply the expansion theorem with $\zeta = \eta/4$, obtaining density α on an interval with $\dim A_0 + \operatorname{codim}_E B_0 \leq r_s$. Intersecting this interval with $\mathcal{C}$ only increases the relative density of $\tilde{S}$, since $\tilde{S} \subseteq \mathcal{C}$, and the intersection is nonempty. A basis (d_i, s_i) of A_0 specifies $A_0 \subseteq L_M$ by the equations $d_i^\top M = s_i^\top$. For the standard dot product on E , a basis (t_j, q_j) of $B_0^\perp$ specifies $L_M \subseteq B_0$ by $Mq_j = t_j$. Their simultaneous solution set is a slice of the required form, with at most r_s equations of either kind. On at least an α fraction of this slice, f equals the value defining S . This proves the conclusion with $z = 0$ and u equal to that value.

The constants α, r_s were chosen using only η . Since $\dim E = \ell + m$, after fixing ℓ the remaining dimension threshold is absorbed into $m_0(\eta, \ell)$. No folding condition was used. The fiber argument is the equality consequence in [6, Lemma 2.3]; numbered references to that paper use its arXiv version 1. $\square$

The constants and dimension thresholds in Theorem 5.1 can be chosen effectively, with rational α when η is rational. This follows from the quantitative bound in [25, Theorem 2.13 and Section 2.4] and the comparison estimates in [25, Lemmas 2.7–2.8], with theorem numbers again referring to the 2018 revision.

We now fix the parameters needed below. Choose $p_* > 0$ and the associated r_* from Lemma 3.1. Set

$$\eta = 2^{-r_*}/16, \quad 2^{-(\ell-r_s)} < \alpha/8, \quad 2^{-\ell} < \eta/4, \quad (5.4)$$

where α, r_s first come from Theorem 5.1, and ℓ is then chosen sufficiently large for that theorem, for the latent lemma, and for the displayed inequalities. Obtain the latent data $(\mathcal{V}, K, C, \mu)$ at this value of ℓ , and fix a linear splitting $\mathcal{V} = K \oplus W$. All these objects are fixed before k grows. In particular, $\ell > r_s$. We write a table as $P = (M, T)$, with $M : L_U \rightarrow K$ and $T : L_U \rightarrow W$.

5.2 The advice experiment and the decoding statement

We first specify the local inputs for two strategies whose answer agreement will contradict the outer-game gap. This experiment is defined for every positive integer k and every $0 < \beta < 1$. Let $\mathcal{Q}_k$ consist of all equation questions $U \in \mathcal{E}^k$ and all permitted projected questions O obtained from them by the coordinate projections of Section 4. The questions U, O and the projection $\bar{\pi} : L_U \rightarrow L_O$ have the law in Section 4: independently in each block, an equation is retained with probability $1 - \beta$, and a uniformly chosen variable position is retained with probability β .

Definition 5.2 (Projected advice experiment). Sample U, O, π , then an independent uniform map $A : K \rightarrow \mathbb{F}_2^{r_s}$ and independent uniform maps

$$T_O : L_O \rightarrow W, \quad M_O : L_O \rightarrow K.$$

Put $S_O = AM_O$ and send the following inputs to the two provers:

$$\begin{aligned} U\text{-prover: } & (U, A, T_O \bar{\pi}, S_O \bar{\pi}), \\ O\text{-prover: } & (O, A, T_O, S_O). \end{aligned} \quad (5.5)$$

The map M_O beyond the stated advice is hidden. Each prover returns an actual answer, that is, a vector of first bit one in its homogeneous answer space. They succeed when their answers a_U, a_O satisfy $\bar{\pi} a_U = a_O$.

The U -message contains the maps $(T_U, S_U) = (T_O \bar{\pi}, S_O \bar{\pi})$, obtained by padding maps sampled on O . We will first find useful advice under a different law: an unrestricted uniform table (M, T) on U , with $S_U = AM$. Later we compare the full padded table $(M_O \bar{\pi}, T_O \bar{\pi})$ with this unrestricted table after averaging over the projection choices. This comparison is the bridge from the extraction argument to the experiment.

Proposition 5.3 (Matrix decoding). *Choose the latent and shortcode parameters in the order specified by Equation (5.4), fixing $\mathcal{V} = K \oplus W$ before k . There are a constant $\gamma > 0$ and an integer $k_0 \geq 8$, depending only on these fixed parameters, with the following property.*

For every cube $k \geq k_0$ and every parity instance, suppose a function

$$F_Q : \text{Hom}(L_Q, \mathcal{V}) \longrightarrow K$$

is given for every $Q \in \mathcal{Q}_k$. Assume folding for every such Q , every $P \in \text{Hom}(L_Q, \mathcal{V})$, and every $c \in K$, and assume sharing for every permitted projection $\bar{\pi} : L_U \rightarrow L_O$ and every $P' \in \text{Hom}(L_O, \mathcal{V})$:

$$\begin{aligned} F_Q(P + ce_0^\top) &= F_Q(P) + c, \\ F_U(P' \bar{\pi}) &= F_O(P'). \end{aligned} \quad (5.6)$$

If this family passes the table test of Equation (4.6) with probability at least 0.99, then the projected advice experiment with $\beta = k^{-2/3}$ has local strategies whose probability of full answer agreement is at least γ .

Every labeling of the vertices of $\mathcal{G}$ supplies such a family: its restored functions satisfy Equation (5.6) by Equations (4.4) and (4.5). The abstract formulation records that no other property of the vertex construction enters the decoding argument. The family is fixed when the two response rules are defined and may be built into those rules.

The next two subsections work for any positive integer k , with a fixed family satisfying Equation (5.6). They use unrestricted uniform tables on U and do not yet restrict β . We specialize k and β when comparing that law with the projected experiment.

5.3 Comparing the two matrix perturbations

The uniform matrix perturbation has the rank-indexed Fourier spectrum computed in [25, Lemma 2.10 in the 2018 revision]. We compare that spectrum with the noise supplied by the latent gadget.

Lemma 5.4 (Spectral comparison). *If the table family passes Equation (4.6) with probability at least 0.99, then*

$$\mathbb{E}_{U,T} \Pr_{M,a,l} [F_U(M,T) = F_U(M + al^\top, T)] \geq 0.9 2^{-r_*}, \quad (5.7)$$

where a is uniform in K and the other matrix and vector samples are uniform in their respective spaces. Consequently at least 10η of the pairs (U,T) have the inner probability at least 4η .

Proof. For $\sigma \in K^*$ define the sign function $f_{U,\sigma}(P) = (-1)^{\sigma(F_U(P))}$. A frequency on the matrix domain is a tuple $g = (g_i)_{i=0}^{m-1}$ of functionals in $\mathcal{V}^*$; its character at P is $(-1)^{\sum_i g_i(Pe_i)}$. All Fourier coefficients use uniform probability measure. Translation by $al^\top$ multiplies this character by $(-1)^{\sum_i l_i g_i(a)}$. Averaging over uniform l therefore gives the eigenvalue

$$\Lambda_g = \Pr_{a \sim \mu} [g_i(a) = 0 \text{ for every } i].$$

It lies in $[0, 1]$. By Lemma 3.1, it is at most $7/8$ whenever

$$\dim \text{span}\{g_i|_K : 0 \leq i < m\} \geq r_*.$$

Indeed the span of the g_i is a subspace of $\mathcal{V}^*$ to which the lemma applies.

Expanding equality by a uniform output character, and then using Fourier orthogonality, expresses the test acceptance as

$$\mathbb{E}_{U,\sigma} \sum_g |\hat{f}_{U,\sigma}(g)|^2 \Lambda_g.$$

Parseval gives total squared mass one for each sign function. If B is the average squared mass on restricted ranks at least r_* , then $0.99 \leq 1 - B/8$, so $B \leq 0.08$. In particular, at least 0.9 of the average mass lies on ranks below r_* .

When a is instead uniform in K , the eigenvalue at g is $2^{-\dim \text{span}\{g_i|_K\}}$. The preceding mass thus gives the lower bound in Equation (5.7). This perturbation leaves T fixed, so the expression there is precisely the average of the ordinary shortcode equality tests at fixed (U, T) . Its lower bound is 14.4η . If q is the probability that the inner acceptance is at least 4η , then the average is at most $4\eta + q$. Thus $q \geq 10.4\eta \geq 10\eta$. $\square$

5.4 Obtaining many useful row fibers

The inverse theorem gives one structured slice on each qualifying (U, T) fiber. The U -prover's information about the K -valued matrix is the pair (A, AM) . We need a positive probability that these observed rows determine a useful slice. We will also identify the short set of valid answers supplied by that slice. The erasure argument below proves the probability bound, and folding then makes the answer set nonempty.

For this extraction step, sample U as before and take independent uniform maps $T : L_U \rightarrow W$, $A : K \rightarrow \mathbb{F}_2^{r_s}$, and $M : L_U \rightarrow K$. For fixed U, T, A and a map $S_0 : L_U \rightarrow \mathbb{F}_2^{r_s}$, the set $\{M : AM = S_0\}$ is its *row fiber*. We call the row advice (U, T, A, S_0) *good* if there exist at most r_s column specifications and a target $Mz + u$ such that the simultaneous solution set of the column specifications and $AM = S_0$ is nonempty and $F_U(M, T) = Mz + u$ with probability at least $\alpha/2$ under the uniform measure on this set.

This is an analytic condition on the row fiber. The target vector z need not initially be an actual answer. The fixed column values allow z to vary modulo the span of their column directions, with a corresponding change in u . Thus the slice naturally supplies a coset of possible target vectors. Its valid part will be the answer set sampled by the U -prover.

Lemma 5.5 (Good row advice and its short answer set). *For all sufficiently large k , a table family with matrix-test acceptance at least 0.99 satisfies*

$$\Pr_{U, T, A, M}[(U, T, A, AM) \text{ is good}] \geq g_0, \quad g_0 = 10\eta^2 2^{-\ell r_s} > 0. \quad (5.8)$$

The lower bound and the required largeness of k are independent of the parity instance and the family.

For every positive integer k and every family satisfying Equation (5.6), one can select deterministically for each good row advice at most r_s column specifications $Mq_j = t_j$, with $q_j \in L_U$ and $t_j \in K$, and a target $Mz + u$, with $z \in L_U$, $u \in K$, and $e_0^\top z = 1$, satisfying the definition of goodness. Writing $Z = \text{span}\{q_j\}$, the chosen witness determines the answer set

$$\mathcal{A} = (z + Z) \cap \{v \in L_U : e_0^\top v = 1\}, \quad 1 \leq |\mathcal{A}| \leq 2^{r_s}. \quad (5.9)$$

The witness and $\mathcal{A}$ depend only on (U, T, A, S_0) and the fixed table family. The slice and its target retain the agreement guarantee in the definition of goodness.

Proof. Fix (U, T) for which the surrogate acceptance is at least 4η , and write $f(M) = F_U(M, T)$. Let B be the union, in the matrix domain, of all good row fibers, over every possible A, S_0 . Suppose that its uniform measure is less than η . Change f at the entries of B , replacing each by an independent uniform element of K . Each matrix receives one replacement, even when several good row fibers overlap there. Both endpoints of the shortcode step are uniform matrices. Consequently every such modification changes acceptance by at most $2 \Pr[B]$, and the modified function has acceptance at least 2η .

We show that the random replacements can be chosen so that they contribute less than $\alpha/4$ of the size of every simultaneous row and column slice as matches with every target $Mz + u$. Put $r = r_s$. A nonempty slice with at most r row and r column specifications has homogeneous space

$$\text{Hom}(\mathbb{F}_2^m / \text{span}\{q_j\}, \bigcap_i \ker d_i^\top).$$

Its size n_S is therefore at least $2^{(\ell-r)(m-r)}$. Padding the lists of specifications with zero equations, the number of slice–target descriptions is at most $2^{(2r+1)(\ell+m)}$.

For a fixed description, count only matches at randomized entries and assign a deterministic zero contribution to the other entries. These are independent variables in $[0, 1]$, with total mean at most $2^{-\ell}n_S < \alpha n_S/8$. Hoeffding's inequality [19, Theorem 2, equation (2.6)] bounds the probability of at least $\alpha n_S/4$ such matches by $\exp(-\alpha^2 n_S/32)$. The union bound over all descriptions is thus at most

$$\exp\left((2r+1)(\ell+m)\log 2 - \frac{\alpha^2}{32}2^{(\ell-r)(m-r)}\right),$$

which tends to zero as $m = 1 + 2k$ grows. Fix replacements with the asserted property.

Consider any slice of the allowed type. If its row fiber was good, every entry of that fiber belongs to B , and its modified agreement is less than $\alpha/4$. If the row fiber was not good, the original agreement with each target on the slice was less than $\alpha/2$. The unmodified matches form a subset of those original matches, and the randomized matches add less than $\alpha/4$. The modified agreement is then less than $3\alpha/4$. Padding shorter row systems by zeros shows that these two cases cover every slice in Theorem 5.1. That theorem now contradicts the modified function's acceptance at least 2η . Its applicability does not require the modification to preserve folding.

It follows that $\Pr[B] \geq \eta$. For each $M \in B$, at least one matrix A has a good row fiber containing M . The uniform choice of A hits that particular matrix with probability $2^{-\ell r_s}$. Hence, at this fixed (U, T) , the probability of good data is at least $\eta 2^{-\ell r_s}$. Lemma 5.4 supplies at least 10η of such pairs (U, T) , proving Equation (5.8). All counting and concentration bounds depended only on the fixed parameters and m , as required.

The remaining assertion does not use high test acceptance or largeness of k . Fix any positive integer k , a family satisfying Equation (5.6), and any good row advice. Fix a witness, write its column specifications as $Mq_j = t_j$, and put $Z = \text{span}\{q_j\}$. Nonemptiness of the slice makes the prescribed values consistent: they determine a linear map $t : Z \rightarrow K$ with $t(w) = Mw$ throughout the slice. Replacing z by $z + w$ and u by $u + t(w)$ for $w \in Z$ therefore preserves the target. It suffices to show that the coset $z + Z$ contains a vector of first bit one.

Otherwise z and every vector in Z have first bit zero. Put $H = \ker A$. For each $h \in H$, translation $M \mapsto M + he_0^\top$ preserves the row fiber, the column specifications, and the target. By folding it translates $F_U(M, T)$ by h . The action is free, and in each of its H -orbits at most one matrix can satisfy the target equality. The agreement is at most

$$|H|^{-1} = 2^{-\ell + \text{rank } A} \leq 2^{-(\ell - r_s)} < \alpha/8,$$

contradicting goodness. Choose $w \in Z$ for which $z + w$ has first bit one and make the corresponding changes to z and u . They preserve the original agreement fraction. Choose a witness and such a representative using a fixed order on the finite possibilities. This choice uses only the row advice and the fixed table family. Its vector z lies in $\mathcal{A}$, so the set is nonempty; since $\dim Z \leq r_s$, its size is at most 2^{r_s} . $\square$

5.5 Transferring the chosen slice

The comparison with sparsely projected questions follows the smooth outer-game method of Khot–Safta [28, Section 3 and Lemma 3.1], with shared linear advice as in Khot–Minzer–Safta [26, Sections 3.3–3.4]. Lemma 5.5 gives a positive probability of row advice that determines a slice and a short set of valid answers under unrestricted sampling on U . In the actual experiment, the maps are sampled on O and padded to U . We compare the complete padded table with the unrestricted table, including the affine intercept, with projection choices averaged out, and then transfer the chosen slice to O .

From now on take k along integer cubes larger than one and set $\beta = k^{-2/3}$ in Definition 5.2.

The choice $\beta = k^{-2/3}$ serves two purposes: $k\beta^2 \rightarrow 0$ makes the padded table law approach the unrestricted law in Lemma 5.6, while $k\beta \rightarrow \infty$ makes the expected number of singleton coordinates diverge. Once the advice spaces are fixed, each singleton has a uniformly positive probability that every visible advice row has zero coefficient on its variable bit. Lemma 6.2 uses these coordinates to bound agreement.

Lemma 5.6 (The padded matrix is asymptotically uniform). *Conditional on every U , the distribution of*

$$(M, T) = (M_O \bar{\pi}, T_O \bar{\pi})$$

has total variation distance $o(1)$ from the unrestricted uniform distribution on $\text{Hom}(L_U, K) \times \text{Hom}(L_U, W)$, as k tends to infinity along cubes. The bound is uniform in U and the parity instance. The projection choices are averaged out in this comparison.

Proof. Combine the two components into a $\mathcal{V}$ -valued table and put $n = |\mathcal{V}|$. On an unprojected equation block the two slopes are independent uniform elements of $\mathcal{V}$. On a singleton block the padded pair of slopes is an equal mixture of

$$(a, 0), \quad (0, a), \quad (a, a), \quad a \text{ uniform in } \mathcal{V}.$$

Call these pair laws P_0 and Q_0 , respectively. The likelihood ratio dQ_0/dP_0 is n at $(0, 0)$, $n/3$ at the other $3(n-1)$ points on the three displayed lines, and zero elsewhere. Thus

$$\chi^2(Q_0 \| P_0) := \mathbb{E}_{P_0} \left[\left(\frac{dQ_0}{dP_0} - 1 \right)^2 \right] = \frac{n-1}{3}.$$

The one-block mixture $(1-\beta)P_0 + \beta Q_0$ consequently has chi-squared divergence $\beta^2(n-1)/3$ from P_0 . Blocks are sampled independently, including when their question IDs coincide. Multiplying their likelihood ratios gives divergence

$$\left(1 + \frac{n-1}{3} \beta^2 \right)^k - 1 = o(1), \tag{5.10}$$

because n is fixed and $k\beta^2 = k^{-1/3}$ tends to zero. Total variation is at most one half the square root of this expression.

For a projection onto the third bit of equation e , the padded intercept acquires the term $b_e a$, where a is that singleton's slope. The original intercept is independent and uniform in $\mathcal{V}$. Adding the sum of these terms leaves it uniform and independent even after all projection choices and slopes are given. The preceding slope estimate therefore proves the assertion for the entire table. Neither the slope laws nor the bound in Equation (5.10) depend on the equation IDs or their right-hand sides. $\square$

Fix good U -advice for a moment, with its chosen witness and answer set $\mathcal{A}$ from Lemma 5.5. For a permitted projection put $z' = \bar{\pi}z$ and $Z' = \bar{\pi}Z$. Since $\bar{\pi}$ preserves the first coordinate,

$$\bar{\pi}\mathcal{A} = (z' + Z') \cap \{v \in L_O : e_0^\top v = 1\}. \tag{5.11}$$

Indeed, every vector in the coset on the right has a preimage in $z + Z$, and the preimage has the same first coordinate. We will show that the O -prover's private Fourier sample reaches this projected answer set.

The certifying slice does not automatically survive projection. For example, a column condition $Mq = t$ becomes $M_O \bar{\pi} q = t$; it is inconsistent if $\bar{\pi} q = 0$ but $t \neq 0$. More generally, projected column directions may acquire new dependencies. The next argument uses the comparison of the full table laws to find a positive probability of samples where the chosen slice remains nonempty and retains enough agreement.

For this analysis, write $D = (U, O, \pi, A, T_O, S_O)$ for the joint data that determine both prover messages and the projection. These data do not specify M_O beyond the constraint $AM_O = S_O$.

Lemma 5.7 (Transfer of a chosen witness). *Suppose the matrix-test acceptance is at least 0.99. For all sufficiently large cubes k , a set of probability at least*

$$\gamma_1 = \frac{\alpha g_0}{8} 2^{-\ell r_s} > 0 \quad (5.12)$$

of the joint data D in Definition 5.2 has the following properties. The padded U -advice is good, with the chosen answer set $\mathcal{A}$ and its certifying witness. Writing $z' = \bar{\pi} z$, that witness transfers to a nonempty slice

$$AM_O = S_O, \quad M_O \bar{\pi} q_j = t_j,$$

on which

$$\Pr[F_O(M_O, T_O) = M_O z' + u \mid AM_O = S_O, M_O \bar{\pi} q_j = t_j \text{ for all } j] \geq \alpha/4. \quad (5.13)$$

The probability in this equation is taken over a uniform matrix in the specified slice.

Proof. First use unrestricted sampling of U, A, M, T , setting $S_0 = AM$. Let $\mathcal{J}$ be the event that (U, T, A, S_0) is good and M satisfies its chosen column specifications. Let $\mathcal{M}$ be equality with the chosen target, declaring this event false off good data. Within a good row fiber the nonempty column slice imposes at most ℓr_s additional scalar linear equations. Hence Lemma 5.5 gives

$$\Pr[\mathcal{J}] \geq g_0 2^{-\ell r_s}, \quad \Pr[\mathcal{J} \cap \mathcal{M}] \geq (\alpha/2) \Pr[\mathcal{J}].$$

It follows that

$$\Pr[\mathcal{J} \cap \mathcal{M}] - (\alpha/4) \Pr[\mathcal{J}] \geq \frac{\alpha g_0}{4} 2^{-\ell r_s}. \quad (5.14)$$

Both events in this expression depend only on (U, A, M, T) . Adding the independent A to Lemma 5.6 preserves its variation bound. We may therefore compare the two event probabilities with their values in the padded experiment. If the total variation distance is $v_k = o(1)$, the difference in Equation (5.14) changes by at most $(1 + \alpha/4)v_k$. For all sufficiently large k the padded difference is at least γ_1 .

Now condition on D . The remaining M_O is uniform on the row fiber $AM_O = S_O$. The padded U -advice, and therefore its chosen witness and answer set, are fixed. On good data let p_D be the conditional probability of the projected column slice, and, when $p_D > 0$, let q_D be its conditional equality fraction. Put $p_D = 0$ on data that are not good, and put $q_D = 0$ when $p_D = 0$. Sharing gives $F_U(M_O \bar{\pi}, T_O \bar{\pi}) = F_O(M_O, T_O)$, so the padded difference is exactly

$$\mathbb{E}_D[p_D(q_D - \alpha/4)] \geq \gamma_1.$$

The integrand is nonpositive outside the event consisting of goodness, $p_D > 0$, and $q_D \geq \alpha/4$, and it is at most one everywhere. That event thus has probability at least γ_1 and has precisely the asserted properties. The law comparison applied to events depending only on (U, A, M, T) ; conditioning on D here only identifies where their positive difference occurs in the projected experiment. $\square$

5.6 Decoding the row advice

Extracting an outer strategy from Fourier mass follows the approach of Håstad [18, Section 1.1] and the Fourier-list method of Khot–Safra [28, Section 1.1]. Here the projected prover samples squared Fourier mass on its own row fiber. The transferred slice will show that this private sample reaches the projected answer set $\bar{\pi}\mathcal{A}$ with positive probability.

We now define both strategies from their separate messages. The joint data and the transferred witness serve only in the analysis of their agreement; the O -prover does not use the U -prover's witness or answer set.

The U -prover uses its visible row advice to make the deterministic choice in Lemma 5.5. When the advice is good, it samples a_U uniformly from $\mathcal{A}$. This is a nonempty set of at most 2^{r_s} actual answers. On other inputs the prover returns any actual answer.

The O -prover starts from its own row fiber. Given (O, A, T_O, S_O) , put $H = \ker A$ and choose, by a fixed rule, a representative M_0 satisfying $AM_0 = S_O$. Such a representative exists on every input in the experiment. Write the row fiber as $M_0 + X$, where $X = \text{Hom}(L_O, H)$. Sample uniform $\sigma \in K^*$ and form the sign function

$$f_\sigma(N) = (-1)^{\sigma(F_O(M_0+N, T_O))}, \quad N \in X.$$

Sample a Fourier frequency $\Phi \in X^*$ with probability $|\hat{f}_\sigma(\Phi)|^2$. Identify $X^* = H^* \otimes L_O$, with the pairing defined on pure tensors by

$$\langle \varphi \otimes v, N \rangle = \varphi(Nv).$$

For $h \in H$, contraction sends Φ to a vector $\Phi(h) \in L_O$. If $\sigma|_H \neq 0$, choose by a fixed rule an $h \in H$ with $\sigma(h) = 1$ and return $\Phi(h)$. If the restriction is trivial, return any actual answer.

This rule always returns an actual answer at frequencies of positive mass. Indeed the folding identity in Equation (5.6) gives

$$f_\sigma(N + he_0^\top) = -f_\sigma(N) \quad \text{when } \sigma(h) = 1.$$

Changing variables in a Fourier coefficient shows that its being nonzero requires $(-1)^{\langle \Phi, he_0^\top \rangle} = -1$, equivalently $e_0^\top \Phi(h) = 1$.

Figure 3 summarizes the information and private choices of the two decoders. The transferred slice will be used only to analyze their agreement probability.




Figure 3: The two decoders use independent random tapes. Neither receives the sampled M_O or the projection $\bar{\pi}$; the right decoder chooses its own representative M_0 from the visible advice. Every displayed advice map includes its affine intercept, through the homogeneous coordinate. The transferred slice is used only to analyze these private rules: the right decoder does not use the left witness or answer set. If the left advice is not good, or if $\sigma|_H = 0$ on the right, that decoder returns an arbitrary valid answer.

Lemma 5.8 (A transferred slice gives agreement). *Fix joint data satisfying Lemma 5.7. Conditional on these data, the two rules just defined, using independent randomness, agree under $\bar{\pi}$ with probability at least*

$$2^{-\ell r_s} (\alpha/8)^2 2^{-r_s}. \quad (5.15)$$

Proof. Put $Z' = \bar{\pi}Z$, so $\dim Z' \leq r_s$. The projected column specifications in the row fiber fix the restriction $N|_{Z'}$ to a linear map $b : Z' \rightarrow H$. The resulting affine slice is nonempty; choose one N_* in it. For every $\sigma \in K^*$ define

$$C_\sigma = \mathbb{E}[f_\sigma(N)(-1)^{\sigma(Nz')} \mid N|_{Z'} = b].$$

The character expansion of Equation (5.13) gives

$$\mathbb{E}_\sigma[(-1)^{\sigma(M_0 z' + u)} C_\sigma] \geq \alpha/4, \quad \text{so} \quad \mathbb{E}_\sigma|C_\sigma| \geq \alpha/4.$$

Exactly a fraction $2^{-\dim H}$ of the characters have trivial restriction to H . Since $\dim H \geq \ell - r_s$ and $|C_\sigma| \leq 1$, Equation (5.4) implies

$$\mathbb{E}_\sigma[\mathbf{1}_{\{\sigma|_H \neq 0\}} |C_\sigma|] \geq \alpha/8. \quad (5.16)$$

The annihilator in X^* of matrices vanishing on Z' is $H^* \otimes Z'$. For $\theta_\sigma = (\sigma|_H) \otimes z'$, Fourier expansion on the affine slice therefore gives

$$C_\sigma = \sum_{\Phi \in \theta_\sigma + H^* \otimes Z'} \hat{f}_\sigma(\Phi) (-1)^{\langle \Phi + \theta_\sigma, N_* \rangle}.$$

This sum has $2^{\dim H \dim Z'} \leq 2^{\ell r_s}$ terms. Cauchy–Schwarz, first on each sum and then in the average over σ , shows from Equation (5.16) that

$$\begin{aligned} \mathbb{E}_\sigma \left[\mathbf{1}_{\{\sigma|_H \neq 0\}} \sum_{\Phi \in \theta_\sigma + H^* \otimes Z'} |\hat{f}_\sigma(\Phi)|^2 \right] \\ \geq 2^{-\ell r_s} \mathbb{E}_\sigma[\mathbf{1}_{\{\sigma|_H \neq 0\}} C_\sigma^2] \geq 2^{-\ell r_s} (\alpha/8)^2. \end{aligned}$$

This is a lower bound on the probability that the O -decoder samples a nontrivial restriction and a frequency in the indicated coset.

Every such frequency has the form $\Phi = (\sigma|_H) \otimes z' + \Psi$ with $\Psi \in H^* \otimes Z'$. For the decoder's choice of h it satisfies

$$\Phi(h) = z' + \Psi(h) \in z' + Z'.$$

As already proved, a frequency of positive mass gives first bit one. Equation (5.11) therefore places the returned vector in $\bar{\pi}\mathcal{A}$. Independently, the U -decoder samples uniformly from $\mathcal{A}$ and chooses a preimage with probability at least $1/|\mathcal{A}| \geq 2^{-r_s}$. Multiplying the bounds proves Equation (5.15). $\square$

Proof of Proposition 5.3. Use the two decoders above. Lemma 5.7 supplies a set of joint data of probability at least γ_1 , and Lemma 5.8 gives the stated conditional agreement on that set. Thus one may take

$$\gamma = \gamma_1 2^{-\ell r_s} (\alpha/8)^2 2^{-r_s} > 0.$$

The largeness requirements for k come only from Theorem 5.1, the concentration bound in Lemma 5.5, and the total variation estimate in Lemma 5.6. They are uniform in the parity instance and table family. Witness choices and Fourier sampling need not be efficient: they define strategies in a finite game for the soundness argument. Finally, averaging over the provers' independent random tapes also gives deterministic strategies with at least the same agreement probability. $\square$

6 Clean coordinates and the soundness gap

Proposition 5.3 converts high acceptance of the matrix test into agreement between two local strategies supplied with correlated linear advice. We now bound the value of this advice experiment on a NO parity instance. The use of vanishing advice is part of the smooth outer-game framework of [26, Sections 3.3–3.4]. We give the full local reconstruction argument for the present occurrence-indexed questions. The key is to identify singleton coordinates on which both advice slopes vanish. After fixing the data outside those coordinates, their question pairs have the ordinary independent equation–variable law. Each side can then reconstruct its original advice input from its own remaining questions and the fixed data. This defines strategies for an ordinary parallel repetition, whose value bounds the original agreement.

6.1 Classical projection-game repetition

We use the following form of the parallel repetition theorem of Dinur and Steurer. The value here is the ordinary classical value: each prover answers using only its own question. A projection constraint maps each answer on one side to its required answer on the other side.

Theorem 6.1 (Projection-game repetition). *Let G be a finite bipartite projection game, with an arbitrary probability distribution on its constraint occurrences. Question sets and answer alphabets may be arbitrary finite sets, and parallel constraint occurrences are allowed. The game $G^{\otimes n}$ samples n independent occurrences, sends each prover its tuple of questions, and accepts if all n projection constraints hold. For every integer $n \geq 1$ and every $0 < g < 1$,*

$$\text{val}(G) \leq 1 - g \implies \text{val}(G^{\otimes n}) \leq (1 - g^2/16)^n. \quad (6.1)$$

The bound is independent of the question-set and alphabet sizes.

This is Corollary 1.2 of [13]. Their Section 2.1 permits nonnegative edge weights and parallel edges, and Section 2.2 defines the independent product used here. Thus the theorem applies to any finite weighted question distribution in our constructions. Interchanging the provers, if necessary, matches the direction of the projection. Private randomization and shared randomness independent of the questions do not increase classical value: sample complete deterministic answer tables using that randomness before the questions arrive, and average their success probabilities.

6.2 Vanishing advice on clean coordinates

Let $\mathcal{E}$ be the parity instance from Section 4, with three distinct variable IDs in each equation occurrence. Recall its ordinary equation–variable game $G_{\mathcal{E}}$: the first prover receives an equation and answers with a satisfying triple; the second receives a uniformly selected variable from that equation and answers with a bit. On a NO instance, Equation (4.1) gives $\text{val}(G_{\mathcal{E}}) \leq 14/15$.

Recall the advice experiment of Section 5, taking k along integer cubes larger than one and setting $\beta = k^{-2/3}$. An ordered tuple U consists of k independent equation samples. Independently in each block, with probability β we keep a uniformly chosen variable position, and otherwise keep the entire equation. This produces O and the answer projection $\bar{\pi} : L_U \rightarrow L_O$ of Equation (4.2). Independently choose

$$A : K \rightarrow \mathbb{F}_2^r, \quad T_O : L_O \rightarrow W, \quad M_O : L_O \rightarrow K$$

uniformly, and put $S_O = AM_O$. The two inputs, respectively, are

$$(U, A, T_O \bar{\pi}, S_O \bar{\pi}) \quad \text{and} \quad (O, A, T_O, S_O). \quad (6.2)$$

The matrix M_O beyond S_O is not part of either input. Each prover must return an actual answer, equivalently a vector whose first coordinate is one in its respective answer space.

Lemma 6.2 (Clean-coordinate bound). *If $\text{OPT}(\mathcal{E}) \leq 4/5$, every pair of local randomized strategies in the advice experiment satisfies*

$$\Pr[\bar{\pi}(a_U) = a_O] \leq \left(1 - \frac{\beta 2^{-\dim W - r_s}}{3600}\right)^k \leq \exp\left(-\frac{2^{-\dim W - r_s}}{3600} k^{1/3}\right). \quad (6.3)$$

In particular, the upper bound tends to zero uniformly over the parity instance as k tends to infinity along cubes, with the advice spaces fixed.

Proof. Fix an arbitrary pair of local randomized strategies in the advice experiment. We first condition on A and combine the two visible advice maps into

$$D_O = (T_O, S_O) : L_O \longrightarrow E_A, \quad E_A = W \oplus \text{im } A.$$

Conditional on the questions and projection choices, the intercept c_O and all block slope columns of D_O are independent uniform elements of E_A . Indeed, the columns of T_O and M_O are independent and uniform, and applying A to a uniform column gives a uniform element of $\text{im } A$. This coefficient law depends only on whether a block is an equation or a singleton, and not on any equation ID, variable ID, or selected position. In particular, c_O is independent of all block data.

Call coordinate i *clean* if it is projected to a singleton and its one advice slope $d_i \in E_A$ is zero. Conditional on A , the clean indicators are independent, with common probability

$$p_A = \frac{\beta}{|E_A|} = \beta 2^{-\dim W - \text{rank } A}. \quad (6.4)$$

More precisely, for every equation occurrence e and position $j \in [3]$,

$$\Pr[e_i = e, i \text{ is a clean singleton at position } j \mid A] = \omega_e \frac{\beta}{3|E_A|}. \quad (6.5)$$

Thus, conditional on being clean, a coordinate has precisely the ordinary equation–variable distribution.

Let $I \subseteq [k]$ be the clean set. For the analysis, let R consist of A , I , the original intercept c_O , and the following data for every $i \notin I$: its equation occurrence, whether the block was retained or projected, the selected position in the latter case, and its advice slope columns. No witness, decoder output, agreement event, or response randomness is included in R .

For each value $R = r$ of positive probability, regard the clean set I recorded in r as fixed. When I is nonempty, we will use r as constant data when defining two strategies for $G_{\mathcal{E}}^{\otimes |I|}$. Each strategy will reconstruct one original advice input from its own product-game questions, apply the corresponding original response rule, and retain the answers in the clean slots.

Conditional on $R = r$, the pairs

$$(e_i, v_{e_i, j_i}), \quad i \in I,$$

are independent question pairs of $G_{\mathcal{E}}$. Indeed, conditional on A , the coordinate data are independent, and Equation (6.5) shows that conditioning on the clean mask preserves the ordinary law in each clean coordinate. The outside coordinate data and the independent intercept impose no further condition on these pairs.

We now verify the two reconstructions, including the affine intercept shared across the blocks. An equation with right-hand side b_e has actual bits $x, y, b_e t_0 + x + y$. Pulling back a singleton slope d_i thus produces the two slopes

$$(d_i, 0), \quad (0, d_i), \quad (d_i, d_i) \quad (6.6)$$

at positions one, two, and three, respectively. The intercept of $D_U = D_O \bar{\pi}$ is

$$c_U = c_O + \sum_{\substack{i \text{ projected to a singleton} \\ j_i=3}} b_{e_i} d_i. \quad (6.7)$$

For a clean coordinate, every slope in Equation (6.6) is zero and its contribution to Equation (6.7) is zero. All remaining terms in that equation are determined by r .

The equation side receives its occurrences e_i for $i \in I$. Together with r , they specify all of U . Its two advice slopes in each clean block are zero, while r determines every other slope and the full intercept. These coefficients reconstruct both padded advice maps without using any clean selected position or variable question. The variable side receives the variable question for each $i \in I$. Together with r , these specify all of O . Its clean singleton slopes are zero, and r supplies every other slope and the original intercept. These coefficients reconstruct its two advice maps without using any clean equation. Repeated IDs across different blocks cause no additional restriction: questions were sampled independently with replacement, and answers use the local copies specified in Section 4.

These reconstructions define the promised product-game strategies: each side applies its original response rule to the reconstructed input and retains its answers in the slots I , in their original order. Their equation answers are satisfying triples because the original answers are valid. For every realization of the clean questions and response randomness, full agreement in the advice experiment implies acceptance in every clean slot. If $I = \emptyset$, the conditional success probability is at most one. Otherwise,

$$\Pr[\bar{\pi} a_U = a_O \mid R = r] \leq \text{val}(G_{\mathcal{E}}^{\otimes |I|}) \leq \left(1 - \frac{1}{3600}\right)^{|I|}.$$

The first inequality uses the strategies just constructed and the conditional product law. The second uses Theorem 6.1 and Equation (4.1). The original response rules are evaluated only at their original local inputs; r indexes the auxiliary product-game strategies. Thus the final bound holds for every r , including $I = \emptyset$ when its right-hand side is interpreted as one.

Finally, conditional on A , the size of I has distribution $\text{Binomial}(k, p_A)$. Averaging the preceding inequality first over R conditional on A , and then over A , gives

$$\begin{aligned} \Pr[\bar{\pi}(a_U) = a_O] &\leq \mathbb{E}_A \left(1 - \frac{p_A}{3600}\right)^k \\ &\leq \left(1 - \frac{\beta 2^{-\dim W - r_s}}{3600}\right)^k \\ &\leq \exp\left(-\frac{k\beta 2^{-\dim W - r_s}}{3600}\right), \end{aligned}$$

since $\text{rank } A \leq r_s$. Substituting $k\beta = k^{1/3}$ proves Equation (6.3). $\square$

6.3 The matrix-game gap

We can now compare the upper bound for every advice strategy with the strategies extracted from a highly accepting matrix labeling.

Proposition 6.3 (A fixed gap with adjustable completeness). *Fix $p_* > 0$ and choose the latent and shortcode data in the order of Equation (5.4). For every sufficiently large cube k , the matrix game $\mathcal{G}$ constructed from a weighted parity instance $\mathcal{E}$ with distinct positions satisfies*

$$\begin{aligned} \text{OPT}(\mathcal{E}) \geq 1 - \xi &\implies \text{val}(\mathcal{G}) \geq 1 - k\xi - p_*/2, \\ \text{OPT}(\mathcal{E}) \leq 4/5 &\implies \text{val}(\mathcal{G}) \leq 99/100. \end{aligned}$$

The lower threshold on k depends only on the fixed latent and shortcode data, not on $\mathcal{E}$ or ξ .

Proof. Completeness is Lemma 4.3. For soundness, fix a labeling of $\mathcal{G}$. Its restored table functions satisfy the folding and sharing identities of Section 4. Proposition 5.3 supplies a constant $\gamma > 0$ and, for every sufficiently large cube k , gives local advice strategies with agreement at least γ whenever those functions have matrix-test acceptance at least 0.99. For all sufficiently large cubes k , the upper bound in Equation (6.3) is smaller than γ . Lemma 6.2 contradicts the existence of those strategies whenever $\text{OPT}(\mathcal{E}) \leq 4/5$. All thresholds are uniform over the parity instance. $\square$

The order matters: p_* can be made arbitrarily small before the gadget parameters are fixed, and ξ can be made arbitrarily small after k is fixed. Thus the same NO gap is compatible with any prescribed small completeness error.

7 Completion of the reduction

The matrix game has a fixed soundness gap and an arbitrarily small completeness error. We now obtain the form required by Theorem 1.1: unweighted, simple bipartite games with arbitrary prescribed errors. Rounding the occurrence weights and subdividing each occurrence will preserve a fixed gap. One application of alphabet-independent repetition then completes the proof.

7.1 Removing weights and repeated edges

We use the following elementary construction for weighted permutation constraints. At this intermediate stage, a constraint occurrence may be a loop, and several occurrences may join the same vertices.

Lemma 7.1 (Rounding and subdivision). *Let H be a finite permutation-constraint game with $h \geq 1$ positive rational occurrence weights summing to one. For every positive rational τ , put $Q = \lceil h/\tau \rceil$. There is a deterministic construction of an unweighted simple bipartite game B with $4Q$ edges and the same alphabet such that*

$$\left| \text{val}(B) - \left(1 - \frac{1 - \text{val}(H)}{4} \right) \right| \leq \frac{\tau}{4}. \quad (7.1)$$

If the constraints of H are translations on a finite binary vector space, so are those of B . For fixed τ , when H is listed by its vertices, occurrences, permutation tables, and binary rational weights, the construction takes polynomial time in its encoding length.

Proof. Write the occurrence weights as $w_1, \dots, w_h$. Start with $n_e = \lfloor Qw_e \rfloor$, and add one to the $Q - \sum_e \lfloor Qw_e \rfloor$ entries having largest fractional parts, breaking ties in a fixed order. Then

$$n_e \geq 0, \quad \sum_e n_e = Q, \quad |w_e - n_e/Q| \leq 1/Q.$$

The weight differences sum to zero. Consequently, for every labeling, the change in the weight of its satisfied-occurrence set is at most

$$\frac{1}{2} \sum_{e=1}^h |w_e - n_e/Q| \leq \frac{h}{2Q} \leq \tau.$$

Replace each occurrence e by n_e copies, omitting it if $n_e = 0$. The resulting unweighted multigraph H' has Q occurrences and $|\text{val}(H') - \text{val}(H)| \leq \tau$.

Orient each occurrence e of H' from u to v , writing its constraint as $a_v = \rho_e(a_u)$. Replace it by the path

$$u \text{ --- } p_e \text{ --- } q_e \text{ --- } r_e \text{ --- } v,$$

using three fresh interior vertices for *every* occurrence. Place all original vertices and all q_e on the left, and all p_e, r_e on the right. In the direction from u to v , put identity permutations on the first three edges and ρ_e on the last. When expressing an edge in left-to-right orientation, invert its permutation if needed. Give all $4Q$ edges equal weight.

Fix labels at the original vertices. A satisfied occurrence extends to four satisfied edges. A violated occurrence cannot extend to four, but propagating labels along the first three edges satisfies exactly three. Since the interiors are disjoint, these extensions can be chosen independently. Maximizing over the original labels gives the exact identity

$$\text{val}(B) = 1 - \frac{1 - \text{val}(H')}{4}. \quad (7.2)$$

This proves Equation (7.1).

A loop becomes a four-cycle with three distinct new vertices. Parallel occurrences have disjoint interiors. Thus each left-right vertex pair carries at most one edge, so B is simple. An inverse binary translation is the same translation, and an identity is translation by zero; the claimed constraint form is preserved.

Finally, $Q \leq h/\tau + 1$ is polynomially bounded for fixed τ . Integer division and exact comparison of rational remainders compute the n_e in polynomial time and bit complexity. Listing their Q copies and the $3Q$ new vertices is therefore polynomial as well. $\square$

7.2 The final reduction

Proof of Theorem 1.1. Fix $\varepsilon, \delta \in (0, 1/2)$. Choose positive rational $\varepsilon_0 \leq \varepsilon$ and $\delta_0 \leq \delta$. We will take $\tau \leq 1/200$ when applying the rounding and subdivision lemma. For any matrix game H satisfying the NO bound $\text{val}(H) \leq 99/100$, the resulting game B then satisfies

$$\text{val}(B) \leq 1 - \frac{1/100 - \tau}{4} \leq 1 - \frac{1}{800}. \quad (7.3)$$

This gap does not depend on the alphabet. We can therefore use alphabet-independent repetition to choose its exponent before the latent construction fixes that alphabet. Choose an integer $t \geq 1$ such that

$$\left(1 - \frac{1}{10240000}\right)^t \leq \delta_0. \quad (7.4)$$

For example, $t \geq \lceil 10240000 \log(1/\delta_0) \rceil$ suffices. By Theorem 6.1, this exponent reduces every projection game of value at most $1 - 1/800$ to value at most δ_0 .

Choose positive rational parameters

$$p_* \leq \frac{\varepsilon_0}{t}, \quad \tau \leq \min \left\{ \frac{1}{200}, \frac{\varepsilon_0}{t} \right\}. \quad (7.5)$$

Choose the latent rank threshold r_* from Lemma 3.1, put $\eta = 2^{-r_*}/16$, and choose the corresponding shortcode parameters α, r_s , using a positive rational lower bound for α . Then choose a sufficiently large ℓ , as specified in Equation (5.4), and fix the corresponding latent data $(\mathcal{V}, K, C, \mu)$ and a splitting $\mathcal{V} = K \oplus W$, with $K = \mathbb{F}_2^\ell$.

Once the rational budgets are fixed, these choices are effective. For rational p_* , the effective assertion of Lemma 3.1 supplies all the finite latent data, including a computable rational noise law, by a terminating deterministic construction. The effective bounds stated after Theorem 5.1 supply the rational choice of α and the dimension thresholds, including $m_0(\eta, \ell)$ after ℓ has been fixed.

Next choose a cube $k \geq 8$ large enough for Proposition 6.3. Such a cube can be found by checking the estimates used in its proof. Write $m = 1 + 2k$ and $\beta = k^{-2/3}$; since k is a cube, β is rational. Require $m \geq m_0(\eta, \ell)$ and

$$\frac{\alpha^2}{32} 2^{(\ell-r_s)(m-r_s)} > (2r_s + 1)(\ell + m).$$

The latter is a sufficient rational condition for the union bound in Lemma 5.5 to be less than one. For witness transfer in Lemma 5.7, require

$$\frac{1}{4} \left[\left(1 + \frac{|\mathcal{V}| - 1}{3} \beta^2 \right)^k - 1 \right] < \frac{\gamma_1^2}{(1 + \alpha/4)^2},$$

where γ_1 is from Equation (5.12). The left-hand side is the square of the total-variation bound. For the final contradiction, require

$$\left(1 - \frac{\beta 2^{-\dim W - r_s}}{3600} \right)^k < \gamma,$$

with γ from Proposition 5.3. All these conditions are rational and hold for sufficiently large cubes k . Enumerating cubes until they hold therefore finds a suitable k ; the search cost depends only on the two prescribed errors and is incurred before the input formula is used.

Finally, with k fixed, choose a positive rational

$$\xi < \min \left\{ \frac{1}{100}, \frac{\varepsilon_0}{2kt} \right\}. \quad (7.6)$$

Every choice depends only on the two prescribed errors. In particular, r_* is fixed before ℓ , the advice dimensions before k , and k before the parity error ξ .

Given a 3SAT formula, apply Theorem 4.1 and Lemma 4.2, and construct the matrix game H of Section 4. Proposition 6.3 gives

$$\begin{aligned} \text{YES: } \text{val}(H) &\geq 1 - d, \quad d = k\xi + p_*/2 \leq \varepsilon_0/t, \\ \text{NO: } \text{val}(H) &\leq 99/100. \end{aligned}$$

Apply Lemma 7.1 with τ to obtain the unweighted simple bipartite game B . It satisfies

$$\begin{aligned} \text{YES: } \text{val}(B) &\geq 1 - (d + \tau)/4, \\ \text{NO: } \text{val}(B) &\leq 1 - 1/800, \end{aligned}$$

where the NO bound is Equation (7.3). The final output is $G = B^{\otimes t}$. Its NO value is at most $\delta_0 \leq \delta$, by Equations (6.1) and (7.4). On a YES input, use a labeling of B coordinatewise. A union bound gives

$$\text{val}(G) \geq 1 - \frac{t(d + \tau)}{4} \geq 1 - \varepsilon_0/2 \geq 1 - \varepsilon.$$

The output has the claimed graph and constraint form. Its two vertex sets are the t -fold products of the two sides of B . A pair of product vertices determines at most one edge in each coordinate, since B is simple, and hence at most one product edge. Independent uniform edge sampling is uniform on these product edges. The constraints of H are translations on K , by Section 4; rounding and subdivision preserve them. The product of translations with offsets $c_1, \dots, c_t$ is translation by $(c_1, \dots, c_t)$ on $K^t = \mathbb{F}_2^{\ell t}$. This is the alphabet $\mathbb{F}_2^s$ in the theorem, with $s = \ell t \geq 1$ fixed by the prescribed errors.

Polynomial encoding. Suppose the preprocessed parity instance has $N_{\mathcal{E}}$ positive equation occurrences. There are $N_{\mathcal{E}}^k$ equation tuples. The number of tables, noise outcomes, and perturbation vectors for each tuple is fixed after the parameter choices above, so the number h of listed matrix-test occurrences is a fixed constant times $N_{\mathcal{E}}^k$. Omit zero-probability outcomes and retain repeated occurrences separately. Each weight is a product of k parity weights, a fixed rational noise probability, and fixed uniform-sampling factors. It has polynomial binary length.

Exact table keys and folded representatives are computed by linear algebra in fixed dimensions, keeping the original variable and occurrence IDs. At most $2h$ endpoints appear. Even comparing every pair of keys to identify equal vertices takes polynomial time; number the resulting distinct vertices consecutively. Largest-remainder rounding uses $Q = \lceil h/\tau \rceil = O(h)$ copies; subdivision creates $3Q$ new vertices and $4Q$ edges. The final product has $(4Q)^t$ edges and polynomially many vertices, because t is fixed. Full permutation tables on the fixed alphabet $\mathbb{F}_2^{\ell t}$ have constant size, and endpoint indices have $O(t \log(h+1))$ bits. These observations give a deterministic polynomial-time explicit construction throughout and complete the proof. $\square$

8 Approximation consequences

Theorem 1.1 turns established conditional hardness results into NP-hardness results. The underlying conditional reductions and approximation thresholds are due to the cited authors. We give the additional arguments needed for the exact instance models stated below. We also record two results from companion manuscripts. For a maximization problem, ratio r means returning value at least r times the optimum; for a minimization problem, factor c means returning value at most c times the optimum. All factors and tolerances below are fixed independently of input size. In each reduction using Theorem 1.1, they determine the Unique Games errors before its alphabet is chosen. If $P \neq NP$, the resulting hardness statements exclude deterministic polynomial-time approximations at the indicated strict thresholds.

8.1 Constraint satisfaction and Vertex Cover

For MAX-CUT, which maximizes the number of edges crossing a bipartition, Khot, Kindler, Mossel, and O’Donnell’s reduction gives NP-hardness of approximation at every fixed ratio in $(\alpha_{\text{GW}}, 1)$, where the Goemans–Williamson constant is $\alpha_{\text{GW}} \simeq 0.87856$ [24, Theorem 1]. This matches the semidefinite approximation threshold of Goemans and Williamson [16]. The reduction uses the Majority Is Stablest theorem of Mossel, O’Donnell, and Oleszkiewicz [32]. For VERTEX COVER, which minimizes the number of vertices meeting every edge, Khot and Regev’s reduction gives NP-hardness at every fixed factor in $(1, 2)$ [27].

A constraint satisfaction problem (CSP) is specified here by a finite nonempty domain D and a finite family Λ of Boolean predicates on D , of bounded arity. A weighted Λ -instance has a finite variable set and a finite collection of predicates from Λ applied to tuples of variables,

with nonnegative rational weights summing to one. An assignment maps the variables to D ; its value is the total satisfied weight, and OPT is the maximum of this value over all assignments. Raghavendra’s theorem characterizes the optimal approximation for every fixed language of this form, up to arbitrarily small positive error, by a basic semidefinite relaxation [36, Theorem 1.1 and Corollary 1.3]. The Vertex Cover result uses its own reduction. Write SDP_R for the value of Raghavendra’s relaxation, as defined in [37, Section 4.5].

The fixed-instance form of Raghavendra’s theorem and the ordering theorem below use a strengthened bipartite formulation of Unique Games. In its YES case, one labeling satisfies every edge incident to at least a $1 - \zeta$ fraction of left vertices; in its NO case, every labeling satisfies at most a ν fraction of edges. These are uniform fractions of the left vertices and edges, respectively. Lemma A.1 gives this promise on an explicit nonempty simple graph whose common positive left degree can meet any fixed lower bound and be divisible by any fixed positive integer. It follows Khot and Regev’s normalization and sequence construction [27, Lemmas 3.3, 3.4, and 3.6], with an additional restriction that makes the graph simple. Its error parameters and degree are fixed before the alphabet is chosen, so Theorem 1.1 supplies the required premise.

Corollary 8.1 (Raghavendra’s consequence). *Let J be a fixed Λ -instance with $A = \text{OPT}(J) < B = \text{SDP}_R(J)$. For every fixed $0 < \eta < (B - A)/2$, it is NP-hard to distinguish Λ -instances I satisfying*

$$\text{OPT}(I) \geq B - \eta \quad \text{from those satisfying} \quad \text{OPT}(I) \leq A + \eta.$$

Proof. For the thesis’s common-arity convention [37, Definitions 7.2.1–7.2.2], identify D with a fixed $[q]$, and pad every predicate with ignored trailing arguments to a common positive arity, filling the added positions with a fixed variable of J , which exists because $A < B$. This leaves each applied payoff and its dependence set unchanged, hence preserves OPT and the relaxation of [37, Section 4.5]. In that relaxation, each local distribution is over assignments to the variables on which its applied payoff actually depends. When emitting an occurrence, discard only the added trailing positions, retaining the original predicate from Λ and its weight.

Apply Raghavendra’s fixed-instance theorem [37, Theorem 7.9] with tolerance $\eta/2$; see also [36, Theorem 1.1 and Lemma 2.3]. The fixed witness and tolerance determine its strengthened Unique Games errors, so the formulation above and Theorem 1.1 supply its premise, with a fixed alphabet $[M]$. Its verifier instances have optimum greater than $B - \eta/2$ in the YES case and at most $A + \eta/2$ in the NO case.

We pass from the verifier’s probability weights to the stated rational weights. For this fixed invocation, the construction in [37, Sections 7.3.1 and 7.5] gives a local probability law θ on a finite set $\mathcal{S}$, after the alphabet $[M]$ is fixed. A record $s \in \mathcal{S}$ specifies an applied predicate of J , its $a(s)$ ordered oracle-query strings in D^M , and the positions filled by their replies. This law may depend on J, η, M , but no varying input parameter enters it. Rational probability vectors are dense in this finite simplex, so fix $\hat{\theta} \in \mathbb{Q}_{\geq 0}^{\mathcal{S}}$ with

$$\sum_{s \in \mathcal{S}} \hat{\theta}(s) = 1, \quad \delta := \frac{1}{2} \sum_{s \in \mathcal{S}} |\hat{\theta}(s) - \theta(s)| < \frac{\eta}{2}.$$

This rational table is fixed finite data of the reduction; its description length may depend on J, η, M .

Let L_Φ and V_Φ be the left and right sides of the nonempty simple game supplied above, and let $h \geq 1$ be its common left degree. List the variables $V_\Phi \times D^M$ used in Section 7.5 of the cited thesis. Its outer verifier chooses a uniform left vertex u and, independently for each oracle query, chooses a uniform neighbor of u and relabels the query string by that edge’s permutation. For each $s \in \mathcal{S}$, $u \in L_\Phi$, and ordered neighbor tuple in $N(u)^{a(s)}$, emit the Λ -predicate prescribed by s on the

resulting query variables. Fill all positions of any syntactic variable omitted from the test with the first listed output variable; the applied predicate is independent of such variables, so this leaves its value unchanged. Give the occurrence weight

$$\frac{\hat{\theta}(s)}{|L_\Phi| h^{a(s)}}.$$

These weights are nonnegative rationals summing to one. Keeping equal occurrences separate gives $|L_\Phi| \sum_{s \in \mathcal{S}} h^{a(s)}$ occurrences and $|V_\Phi| |D|^M$ variables. The local set and every $a(s)$ are fixed, while the graph has polynomial size; hence this is a polynomial occurrence list. The explicit neighbor lists and edge permutations let us write it in polynomial time. Each weight has polynomial binary length: it can be represented with a fixed denominator multiplied by $|L_\Phi| h^{a(s)}$.

For an assignment σ , let $g_\sigma(s) \in [0, 1]$ be its conditional acceptance probability over these outer choices. The original and rationalized instances use the same variable set, and

$$|\text{val}_{\hat{\theta}}(\sigma) - \text{val}_{\theta}(\sigma)| = \left| \sum_{s \in \mathcal{S}} (\hat{\theta}(s) - \theta(s)) g_\sigma(s) \right| \leq \delta < \frac{\eta}{2}.$$

Taking maxima gives the same bound for their optima. The imported gap therefore yields optimum at least $B - \eta$ in the YES case and at most $A + \eta$ in the NO case, as required. $\square$

8.2 Ordering constraints

An ordering CSP uses a global order of the input variables, rather than an assignment from a fixed finite domain. Fix an integer $k \geq 2$ and a nonempty proper subset $\Pi \subset S_k$, where S_k is the set of permutations of $\{1, \dots, k\}$. In the finite probability-weighted model of Guruswami, Håstad, Manokaran, Raghavendra, and Charikar [17, Definition 8.2], we use an explicitly listed finite variable set V and a finite list of constraint occurrences. Their nonnegative rational weights are encoded by binary numerators and denominators and sum to one. Every occurrence is an ordered k -tuple $(v_1, \dots, v_k)$ of distinct variables, satisfied by a linear order $\prec$ when

$$v_{\pi(1)} \prec \dots \prec v_{\pi(k)} \quad \text{for some } \pi \in \Pi.$$

Write $\text{OPT}_\Pi(I)$ for the maximum satisfied weight. A uniform random order satisfies expected weight $\rho_\Pi = |\Pi|/k!$. Since $\text{OPT}_\Pi(I) \leq 1$, it is a ρ_Π -approximation. For fixed k , conditional expectation derandomizes this guarantee. After any prefix of the order has been exposed, a constraint has at most $k!$ possible relative orders consistent with that prefix. Its conditional success probability is therefore computable exactly; choosing each next variable to preserve the expected value takes a polynomial number of exact rational operations of polynomial bit length. The size of V is part of the input, so this is not a CSP over the fixed finite domain used in the preceding corollary.

Corollary 8.2 (Ordering-CSP consequence). *Fix $k \geq 2$ and $\emptyset \neq \Pi \subsetneq S_k$, and put $\rho_\Pi = |\Pi|/k!$. For every fixed $0 < \varepsilon < (1 - \rho_\Pi)/2$, it is NP-hard to distinguish finite weighted Π -ordering-CSP instances I satisfying*

$$\text{OPT}_\Pi(I) \geq 1 - \varepsilon \quad \text{from those satisfying} \quad \text{OPT}_\Pi(I) \leq \rho_\Pi + \varepsilon.$$

Consequently, approximation at every fixed ratio $r \in (\rho_\Pi, 1]$ is NP-hard. Together with the random-ordering guarantee, this gives the optimal threshold ρ_Π .

Proof. Proposition A.2 constructs the stated finite instance from a simple strong-left Unique Games instance. It uses the local gap and analytic estimates of Guruswami et al. [17] and fixes the required game errors and degree conditions before the alphabet is chosen. Lemma A.1 supplies that game from Theorem 1.1; composing the reductions proves the displayed gap for the original predicate Π .

For a fixed $r \in (\rho_\Pi, 1]$, choose the positive error still smaller so that $\varepsilon < (r - \rho_\Pi)/(1 + r)$. Then $r(1 - \varepsilon) > \rho_\Pi + \varepsilon$, so an r -approximation would distinguish the two cases. $\square$

For MAXIMUM ACYCLIC SUBGRAPH in the finite weighted directed-graph model, take $k = 2$ and $\Pi = \{12\}$ in one-line notation. A constraint (u, v) asks that the edge point forward; forward edges form an acyclic subgraph, and every acyclic subgraph has a topological order. Its threshold is therefore $1/2$. For BETWEENNESS, the constraint (a, b, c) asks that b lie between a and c , so $\Pi = \{123, 321\} \subset S_3$ and the threshold is $2/3! = 1/3$ [17, Theorem 1.1 and Section 1.2]. These finite weighted consequences have near-perfect completeness with fixed positive error.

8.3 Cut, deletion, and clustering consequences

For the cut problems of Chawla, Krauthgamer, Kumar, Rabani, and Sivakumar [9, Section 1], let $G = (V, E)$ be a finite undirected graph with positive rational edge costs c_e , and let $\mathcal{D} = \{\{s_i, t_i\} : 1 \leq i \leq k\}$ contain $k \geq 1$ distinct demand pairs with $s_i \neq t_i$. For $M \subseteq E$, write $c(M) = \sum_{e \in M} c_e$, and let $q_{\mathcal{D}}(M)$ count the demand pairs disconnected in $G - M$. Weighted MULTICUT minimizes $c(M)$ subject to $q_{\mathcal{D}}(M) = k$. Their nonuniform SPARSEST CUT search problem minimizes $c(M)/q_{\mathcal{D}}(M)$ over cutsets M with $q_{\mathcal{D}}(M) > 0$. Only the selected pairs carry unit demand; they need not be all vertex pairs.

For weighted Min-2CNF⁼ Deletion, the input consists of clauses $(\ell \equiv \ell')$ with positive rational weights, where ℓ, ℓ' are Boolean literals, and an assignment pays the total weight of false clauses. The allowed clauses therefore express equality or disequality of variables.

In the arbitrary weighted-graph form of CORRELATION CLUSTERING, each present edge has a $+$ or $-$ label and a nonnegative rational weight. A partition into any number of clusters pays for $+$ edges whose endpoints lie in different clusters and $-$ edges whose endpoints lie in the same cluster. Missing edges carry no penalty. This is the minimum-disagreements objective of [10, Section 2.1].

Corollary 8.3 (Cut, deletion, and clustering hardness). *For every fixed real constant $C > 1$, it is NP-hard to approximate each of the four minimization problems above within factor C . For SPARSEST CUT, this is hardness of the search task under polynomial-time Cook reductions: the algorithm must output a cutset, not only an estimate of the optimum. Consequently, unless $P = NP$, none of these problems admits a deterministic polynomial-time approximation with any fixed constant factor.*

Proof. Corollary 1.3 of [9] gives the first three conclusions under its ordinary edge-value Unique Games hypothesis. For a requested factor C , use its hardness statement at any fixed larger factor L : a C -approximation would also be an L -approximation. To supply that hypothesis, orient the bipartite constraints of Theorem 1.1 left-to-right, renumber its fixed alphabet as $[d]$ with $d = 2^s \geq 2$, and give its edges equal weights summing to one. Padding the smaller vertex side with isolated questions if needed preserves both value and polynomial size. The search-version qualification for SPARSEST CUT is in Section 1.3 of the cited source.

For CORRELATION CLUSTERING, the hard MULTICUT instances in [9, Section 2.1] have distinct antipodal demand pairs inside d -cubes, while the edges of the Multicut graph are cube edges or join different cubes. Here $d \geq 2$, unchanged by Lemma 1.5 of that source, so the demand pairs are nonedges. The transformation of Demaine, Emanuel, Fiat, and Immorlica [10, Theorem 4.7

and Corollary 4.8] therefore creates no opposite-signed parallel edges. It preserves the optimum and converts every clustering into a multicut of no greater cost, so the same fixed-factor hardness transfers. $\square$

The separate manuscript [33, Theorem 1.1] proves NP-hardness at every fixed factor $C > 1$ for the bipartition objective with nonnegative rational capacities and exactly unit demand between every pair of distinct vertices. This uniform-demand theorem does not depend on Theorem 1.1 and is not supplied by the nonuniform reduction above.

Kernel clustering. A separate companion application concerns the following maximization problem. For *identity-target kernel clustering*, fix an integer $k \geq 3$. An input is an explicitly represented rational symmetric matrix $A = (a_{pq})_{p,q=1}^N \succeq 0$ with $A\mathbf{1} = 0$, where $\mathbf{1}$ is the all-ones vector. An assignment $\sigma : \{1, \dots, N\} \rightarrow \{1, \dots, k\}$ has value

$$\text{val}_A(\sigma) = \sum_{p,q=1}^N a_{pq} \mathbf{1}_{\{\sigma(p)=\sigma(q)\}}, \quad \text{OPT}_k(A) = \max_{\sigma} \text{val}_A(\sigma).$$

The sum includes ordered pairs and the diagonal; empty clusters are allowed, and there is no normalization by cluster size. A loss factor $\alpha \geq 1$ means returning value at least $\text{OPT}_k(A)/\alpha$. The Gaussian propeller companion uses Theorem 1.1 to prove NP-hardness of achieving every fixed loss factor

$$1 \leq \alpha < \alpha_k, \quad \alpha_k = \frac{8\pi}{9} \left(1 - \frac{1}{k}\right),$$

for each fixed $k \geq 3$ [35, Theorem 6.1]. This is strict hardness below the propeller rounding factor α_k .

A Finite formulations of the consequences

Some of the reductions in Section 8 use a bipartite Unique Games instance in which a nearly satisfying labeling satisfies all constraints incident to most left vertices. We first obtain a simple graph with this property and with a prescribed lower bound on its common left degree. We then use it to give the finite rational formulation of the ordering consequence. The local ordering construction follows Guruswami, Håstad, Manokaran, Raghavendra, and Charikar [17] and uses the low-influence results cited below. The composition specifies the graph, the independent neighbor choices, and the conversion to constraints on distinct variables.

A.1 A simple strong-left form of Unique Games

Orient the constraints of a bipartite game from its left side to its right side. A left vertex is *strongly satisfied* by a labeling when every edge incident to it is satisfied.

Lemma A.1 (Simple strong-left games). *Fix $0 < \zeta, \nu < 1$ and integers $d, H \geq 1$. There are integers $h, R \geq 1$, with $h \geq H$ and $d \mid h$, and a deterministic polynomial-time reduction from 3SAT to nonempty explicit unweighted simple bipartite Unique Games instances $\Phi_\varphi = (A, B, E)$ over $[R]$ with the following properties. Every left vertex has degree h . If φ is satisfiable, one labeling strongly satisfies at least a $1 - \zeta$ fraction of A . If φ is unsatisfiable, then $\text{val}(\Phi_\varphi) < \nu$. The degree h and the two errors used in Theorem 1.1 may be chosen from ζ, ν, d, H before that theorem supplies the alphabet R .*

Proof. We adapt the normalization and sequence construction of Khot and Regev [27, Lemmas 3.3, 3.4, and 3.6]. The extra copies of right vertices used below ensure that the final graph is simple.

Choose a multiple h of d with

$$h \geq \max\{H, 2\}, \quad \frac{1}{h} < \frac{\nu}{4}.$$

For $(T)_h = T(T-1)\cdots(T-h+1)$, choose a fixed integer $T \geq h$ such that

$$p_{\text{graph}} := \frac{(T)_h}{T^h} > \frac{3}{4}.$$

Such a choice exists because $1 - (T)_h/T^h \leq \binom{h}{2}/T$. Choose positive rationals β, γ, u and a positive integer L_2 so that

$$\beta < \frac{\nu}{4}, \quad h^2\gamma < \frac{\nu}{4}, \quad (h+1)u < \frac{\zeta}{2}, \quad \frac{1}{L_2} < \min\left\{\frac{\beta\gamma}{2}, \frac{\zeta}{2h}\right\}.$$

Finally choose positive rational $\gamma_0, \zeta_0 < 1/2$ with

$$2\gamma_0 < \frac{\beta\gamma}{2}, \quad \sqrt{2\zeta_0} < u.$$

All these choices precede the invocation of Theorem 1.1 with completeness error ζ_0 and soundness error γ_0 .

Write its output as $\Phi_0 = (X_0, Y, E_0)$, orienting each permutation left-to-right and renumbering the fixed alphabet as $[R]$. Put $M_0 = |E_0| > 0$, give each edge weight $1/M_0$, and let w_x be the total incident weight at x . Simplicity gives one permutation at each pair of positive weight. Create

$$n_x = \lfloor 2|X_0|w_x \rfloor$$

copies of x , each with local edge weights w_{xy}/w_x . Only vertices with $w_x > 0$ have copies. If X_1 is the copy set, then

$$|X_0| \leq |X_1| = \sum_x n_x \leq 2|X_0|,$$

and the incident weight at each copy is one.

Here and below a local satisfied weight uses the labels at that left vertex and at the shared right vertices. Given any labeling of the copy game, keep the right labels and give all copies of each x a label that maximizes their common local objective. The corresponding labeling L of Φ_0 satisfies

$$\sum_{x' \in X_1} w_{\text{sat}}(x') \leq \sum_{x: n_x > 0} \frac{n_x}{w_x} w_{\text{sat}, L}(x) \leq 2|X_0| w_{\text{sat}, L}(\Phi_0).$$

Consequently the copy game's normalized value is at most $2\gamma_0$ on a NO input. On a YES input, lift a labeling of Φ_0 with unsatisfied weight at most ζ_0 . The same calculation for unsatisfied weight gives mean local loss at most $2\zeta_0 < u^2$. Markov's inequality leaves more than a $1 - u$ fraction of X_1 with local satisfied weight at least $1 - u$.

To remove the weights, put $\alpha = L_2|Y|$. At each $x \in X_1$, write $\tilde{w}_{xy}$ for its normalized local weights. Choose one neighbor y_0 of positive local weight, round $\alpha\tilde{w}_{xy}$ down for every $y \neq y_0$, and assign the remaining occurrences to y_0 . All occurrences at a pair retain that pair's permutation. This gives an unweighted multigraph Φ_{occ} of left degree α . For any subset of the neighbors, the difference between its old weight and its new occurrence fraction has magnitude at most $|Y|/\alpha = 1/L_2$: sum

the rounding losses over the subset if it omits y_0 , and over its complement otherwise. This applies to the satisfied subset for every labeling. Hence

$$\text{val}(\Phi_{\text{occ}}) \leq 2\gamma_0 + \frac{1}{L_2} < \beta\gamma \quad \text{on a NO input.} \quad (\text{A.1})$$

On a YES input, the labeling above has satisfied occurrence fraction at least $1-a$, where $a = u+1/L_2$, at more than a $1-u$ fraction of X_1 .

Index the α occurrences at x by j , with right endpoint y_j and permutation π_j . Replace each $y \in Y$ by (y, t) , $t \in [T]$, and replace occurrence j by all T occurrences from x to (y_j, t) with permutation π_j . Call the resulting left- αT -regular multigraph Φ_{tag} . For any labeling of Φ_{tag} , keep its left labels and, independently for each original y , give y the label of a uniformly chosen copy (y, t) . The expected edge value in Φ_{occ} , counting its occurrences, is exactly the given edge value in Φ_{tag} . Thus

$$\text{val}(\Phi_{\text{tag}}) \leq \text{val}(\Phi_{\text{occ}}) < \beta\gamma \quad \text{on a NO input.} \quad (\text{A.2})$$

The tag-independent lift of the YES labeling preserves every local satisfied fraction.

Consider the full sequence graph whose left vertices are

$$(x, (j_1, t_1), \dots, (j_h, t_h)), \quad j_i \in [\alpha], \quad t_i \in [T],$$

with position i joined to (y_{j_i}, t_i) by π_{j_i} . This is the sequence construction of Khot and Regev; its left vertex identifier includes the entire ordered occurrence sequence. We record the soundness estimate in a form that permits the left label to depend on that entire sequence. Each x contributes exactly $(\alpha T)^h$ full sequence vertices, all of degree h . Thus the uniform edge value is the average satisfied fraction over a uniform x , an ordered h -tuple of independent uniform occurrences at x , and a uniform position in that tuple.

Fix arbitrary right labels in Φ_{tag} . Each occurrence at x determines the unique left label that satisfies it. Let $p_x(\ell)$ be the fraction determining label ℓ . The $p_x(\ell)$ sum to one, and choosing the best left label at each x gives

$$\mathbb{E}_x \max_{\ell} p_x(\ell) \leq \text{val}(\Phi_{\text{tag}}) < \beta\gamma.$$

Fewer than a β fraction of x therefore have $\max_{\ell} p_x(\ell) \geq \gamma$. At every other x , the probability that two labels determined by h independent occurrences agree is at most

$$\binom{h}{2} \sum_{\ell} p_x(\ell)^2 \leq \binom{h}{2} \max_{\ell} p_x(\ell) < h^2 \gamma.$$

If no two determined labels agree, any label of the sequence vertex satisfies at most one of its h edges. On the remaining sequences its satisfied fraction is at most one. Averaging first over sequences and then over x , and using $1/h + h^2 \gamma < 1$, bounds the full sequence graph's value by

$$s := \beta + \frac{1}{h} + (1 - \beta)h^2 \gamma < \frac{3\nu}{4}. \quad (\text{A.3})$$

The estimate holds for every choice of the right labels and every sequence-dependent choice of left labels.

Retain only sequence vertices for which $t_1, \dots, t_h$ are pairwise distinct. For each x and each base sequence $(j_1, \dots, j_h)$, exactly $(T)_h$ tag sequences remain. Each x thus contributes $\alpha^h (T)_h$ retained vertices, a fraction p_{graph} of its full set. The h right neighbors of each retained vertex are distinct

because their tags are distinct. Different occurrence sequences remain different left vertices, so this retained graph is simple and has common left degree h .

Give each sequence vertex based at x the old label of x , and keep the tag-independent right labels. For this lifted YES labeling, satisfaction at position i depends only on j_i . Conditioning the tags to be distinct leaves the base occurrence sequence independent and uniform. If the satisfied fraction at x in Φ_{occ} is f_x , exactly an f_x^h fraction of its retained sequence vertices are strongly satisfied. Since every x contributes equally, the strongly satisfied fraction is at least

$$(1-u)(1-a)^h \geq 1-u-ha = 1-(h+1)u - \frac{h}{L_2} > 1-\zeta.$$

Here $0 < a < 1$ follows from the parameter choices.

For NO, extend any labeling of the retained graph to the full graph by giving discarded left vertices arbitrary labels and keeping all right labels. All left degrees are h , so its full edge value is at least p_{graph} times its retained edge value. By (A.3), the latter is at most

$$\frac{s}{p_{\text{graph}}} < \frac{3\nu/4}{3/4} = \nu.$$

This extension permits arbitrary dependence of the retained left labels on their tags.

All the operations are finite and deterministic. If N bounds the explicit size of Φ_0 , then $|X_1| = O(N)$ and $\alpha = L_2|Y| = O(N)$. The retained graph has $|X_1|\alpha^h(T)_h = O(N^{h+1})$ left vertices and h edges per left vertex. Its identifiers and the rational arithmetic used for rounding have polynomial bit length, and its permutation tables use the fixed alphabet inherited from Φ_0 . This proves the claimed polynomial bound and completes the reduction. $\square$

A.2 The finite ordering transfer

The ordering theorem of Guruswami et al. uses a multiscale local instance and a low-influence rounding estimate. Their general outer composition is described as analogous to the two-query construction in their Section 7. The following proposition gives the composition for the rational occurrence-list model of Section 8. It separates the constants chosen for the ordering gap from the subsequently chosen Unique Games alphabet.

Proposition A.2 (Finite ordering transfer). *Fix $k \geq 2$ and $\emptyset \neq \Pi \subsetneq S_k$, put $\rho_\Pi = |\Pi|/k!$, and fix $0 < \delta < (1 - \rho_\Pi)/2$. There are integers $q, H \geq 1$ and positive rationals $\zeta, \nu < 1$, depending only on k, Π, δ , with the following property for every fixed integer $R \geq 1$.*

Let $\Phi = (A, B, E)$ be a nonempty explicit unweighted simple bipartite Unique Games instance over $[R]$, oriented from A to B , whose common left degree h satisfies $h \geq H$ and $q \mid h$. There is a deterministic transformation, polynomial in the explicit size of Φ for fixed R , producing an explicit list of Π -ordering constraints on distinct variables within each constraint, with nonnegative rational weights, each encoded by a binary numerator and a positive binary denominator, summing to one. If some labeling strongly satisfies at least $(1 - \zeta)|A|$ left vertices, then $\text{OPT}_\Pi(I) \geq 1 - \delta$. If $\text{val}(\Phi) \leq \nu$, then $\text{OPT}_\Pi(I) \leq \rho_\Pi + \delta$. The constants q, H, ζ, ν are independent of both R and h .

The output variables will be $(b, z) \in B \times [m]^R$, for a fixed local symbol set $[m]$ chosen below. At each $a \in A$, restrict a global order to $N(a) \times [m]^R$ and divide that local sorted list into q equal consecutive blocks. Averaging block membership over neighbors, with their edge permutations acting on z , gives one simplex-valued function of z , whose components have mean $1/q$ under uniform z . For fixed a , conditional on the local query strings, independent neighbor choices make the block ranks

independent with probabilities given by this function. The local estimate bounds their expected coarse payoff under small influences; the collision estimate allows correlated strings and bounds the chance of a shared block, where coarsening can change the induced order. A large influence in an averaged rank interval instead yields a game label, and many such vertices contradict low edge value. We then condition on distinct neighbors to obtain valid constraints and refine the coarse completeness map to a strict order.

For $P = \mathbf{1}_\Pi$, choose $\pi_0 \in \Pi$ and set

$$P_0(\sigma) = P(\pi_0 \circ \sigma), \quad \sigma \in S_k.$$

Then $P_0(\text{id}) = 1$, and its average is ρ_Π . For a vector of integer ranks $r = (r_1, \dots, r_k)$, let $\text{Ref}(r)$ be the permutations compatible with their weak order: they list positions in increasing rank and may order positions of equal rank arbitrarily. Define

$$\overline{P}_0(r) = \frac{1}{|\text{Ref}(r)|} \sum_{\sigma \in \text{Ref}(r)} P_0(\sigma). \quad (\text{A.4})$$

This is the uniform tie extension used in [17, Section 8.2]. It lies in $[0, 1]$ and agrees with P_0 when the ranks are distinct.

The normalization preserves the predicate after one permutation of tuple positions. Given a raw tuple $w = (w_1, \dots, w_k)$, define the emitted tuple v by

$$v_{\pi_0(j)} = w_j \quad (j \in [k]). \quad (\text{A.5})$$

If a strict order induces σ on w , it induces $\pi_0 \circ \sigma$ on v . More generally, relabeling positions gives the bijection

$$\text{Ref}(f(v_1), \dots, f(v_k)) = \{\pi_0 \circ \sigma : \sigma \in \text{Ref}(f(w_1), \dots, f(w_k))\}$$

for every rank map f . Uniform averages over these sets therefore give equal P - and P_0 -payoffs, including ties. The permutation of positions in (A.5) preserves distinctness and weights.

For a finite local instance $J = (U, \mathcal{C}, \omega)$, write $t_c = (u_{c,1}, \dots, u_{c,k})$ for the tuple of occurrence $c \in \mathcal{C}$ and ω_c for its probability weight. Its q -ordering value is

$$\text{opt}_q(J) = \max_{g: U \rightarrow [q]} \sum_{c \in \mathcal{C}} \omega_c \overline{P}_0(g(u_{c,1}), \dots, g(u_{c,k})).$$

Lemma A.3 (Rational local ordering data). *For every integer $q \geq 1$ and rational $\eta > 0$, there are a finite local instance $J = (U, \mathcal{C}, \omega)$, with $|U| = m \geq 2$, and probability laws μ_c on $[m]^k$ with the following properties. Every tuple c has k distinct variables; the weights ω_c and all atoms of μ_c are rational. Every one-position marginal of μ_c is uniform on $[m]$, and there is a fixed $0 < \alpha \leq 1/2$ such that $\mu_c(x) > \alpha$ for every c, x . Moreover*

$$\text{opt}_q(J) \leq \rho_\Pi + \eta, \quad v_{\text{loc}} := \sum_c \omega_c \mathbb{E}_{X \sim \mu_c} \overline{P}_0(X) \geq 1 - \eta. \quad (\text{A.6})$$

The local laws have common feasible Gram data for the relaxation in [17, Section 8.3]: there are a unit vector I and vectors $b_{u,j}$, $u \in U, j \in [m]$, with

$$\sum_j b_{u,j} = I, \quad \langle b_{u,i}, b_{u,j} \rangle = 0 \ (i \neq j), \quad \|b_{u,j}\|^2 = \frac{1}{m},$$

and, for every two positions r, s of c ,

$$\langle b_{u_{c,r},i}, b_{u_{c,s},j} \rangle = \Pr_{X \sim \mu_c} [X_r = i, X_s = j]. \quad (\text{A.7})$$

All these data are fixed by k, Π, q, η .

Proof. The explicit increasing-tuple construction in [17, Theorem 11.1, Definition 11.2, and the proof of Theorem 2.5], applied to P_0 , gives a finite instance J^* with $\text{opt}_q(J^*) \leq \rho_\Pi + \eta$ and a linear order satisfying every constraint. Its distribution uses only uniform choices from finite sets, so its weights are rational. Its tuple positions lie in successive disjoint intervals and are therefore distinct.

Use the finite-copy smoothing from the proof of [17, Lemma 8.9], specialized to this complete order. Take L disjoint copies of J^* , each of weight $1/L$, and place their complete orders consecutively in an order $\sigma : U \rightarrow [m]$, where $m = L|U^*|$. The resulting instance J still has $\text{opt}_q(J) \leq \rho_\Pi + \eta$, since the value of any q -assignment is the average of its values on the copies. Choose L and a positive rational $s_0 < 1$ so that $1/L + s_0 < \eta$. Define a random assignment $X : U \rightarrow [m]$ as follows: with probability $1 - s_0$, choose a uniform cyclic shift of the ranks of σ ; with probability s_0 , assign the variables independently and uniformly from $[m]$.

Each variable is uniform in either part of the mixture. A cyclic cut can split at most one of the L consecutive copies, so every shift has payoff at least $1 - 1/L$. The mixture's payoff is therefore at least $(1 - s_0)(1 - 1/L) > 1 - \eta$. Let μ_c be its restriction to the variables of c . If $n_c(x)$ is the number of the m cyclic shifts with restriction x , then

$$\mu_c(x) = (1 - s_0) \frac{n_c(x)}{m} + \frac{s_0}{m^k}.$$

These atoms are rational and exceed $\alpha := s_0/(2m^k)$. The denominator m^k is valid because the variables in c are distinct.

Finally regard the finite mixture as a probability space and take $I = \mathbf{1}$ and $b_{u,j} = \mathbf{1}_{\{X(u)=j\}}$ in its L^2 space. Uniform marginals, disjoint events for different labels of one variable, and restrictions to each c give all the displayed Gram identities. This constructs the required feasible data for this particular local instance. $\square$

For scalar functions on the uniform product $[m]^R$, use

$$\text{Inf}_\ell(F) = \mathbb{E}[\text{Var}_{z_\ell}(F \mid z_{-\ell})].$$

The noise operator T_θ , $0 \leq \theta \leq 1$, independently keeps each coordinate with probability θ and otherwise resamples it uniformly, then averages the function. In an orthonormal product expansion these are

$$T_\theta F = \sum_{\sigma} \theta^{|\sigma|} \widehat{F}(\sigma) \chi_\sigma, \quad \text{Inf}_\ell(F) = \sum_{\sigma: \sigma_\ell \neq 0} \widehat{F}(\sigma)^2. \quad (\text{A.8})$$

In particular, further smoothing decreases every influence.

Lemma A.4 (Uniform collision estimate). *For every $0 < \kappa < 1/2$, there is $\mu_0(\kappa) > 0$. For each $m \geq 2$, there is a nonnegative function $r_c^{m,\kappa}(t) \rightarrow 0$ as $t \downarrow 0$ with the following property, uniformly over all $R \geq 1$. Let $F, G : [m]^R \rightarrow [0, 1]$ have the same mean $0 < \mu \leq \mu_0(\kappa)$ and satisfy*

$$\text{Inf}_\ell(T_{1-\kappa} F) \leq t, \quad \text{Inf}_\ell(T_{1-\kappa} G) \leq t \quad (\ell \in [R]).$$

For every coupling (Z, Z') whose two marginals are uniform on $[m]^R$,

$$\mathbb{E}[(T_{1-2\kappa} F)(Z)(T_{1-2\kappa} G)(Z')] \leq \mu^{1+\kappa/2} + r_c^{m,\kappa}(t). \quad (\text{A.9})$$

The function $r_c^{m,\kappa}$ is independent of R, μ, F, G and the coupling.

Proof. This is the dimension-uniform form of [17, Lemmas 3.4 and 3.5] needed here. For $H = T_{1-\kappa}F$, the Fourier formula gives

$$\|T_{1-2\kappa}F\|_2^2 \leq \text{Stab}_{1-\kappa}(H),$$

where $\text{Stab}_\theta(H) = \sum_\sigma \theta^{|\sigma|} \widehat{H}(\sigma)^2$. Indeed, $(1-\kappa)^3 - (1-2\kappa)^2 = \kappa(1-\kappa-\kappa^2) > 0$. Each full coordinate influence bound for H implies the corresponding low-degree influence bound in Theorem 4.4 of Mossel, O'Donnell, and Oleszkiewicz [32]. Apply that theorem with product dimension R , one-coordinate atom bound $1/m$, and correlation $1-\kappa$. For all sufficiently small t , it gives

$$\text{Stab}_{1-\kappa}(H) \leq \Gamma_{1-\kappa}(\mu) + C \frac{\log(2m)}{\kappa} \frac{\log \log(1/t)}{\log(1/t)}.$$

The cutoff and this error are independent of R and μ . Here $\Gamma_\theta(\mu)$ is the Gaussian halfspace noise stability of a set of Gaussian measure μ .

For fixed κ , Theorem B.5 of the same source gives $\Gamma_{1-\kappa}(\mu) = \mu^{2/(2-\kappa)+o(1)}$ as $\mu \downarrow 0$. Since

$$\frac{2}{2-\kappa} - \left(1 + \frac{\kappa}{2}\right) = \frac{\kappa^2}{2(2-\kappa)} > 0,$$

there is a cutoff $\mu_0(\kappa) > 0$, independent of m, R , below which $\Gamma_{1-\kappa}(\mu) \leq \mu^{1+\kappa/2}$. This proves the corresponding squared L^2 bound for both F and G , with an error $r_c^{m,\kappa}(t) \rightarrow 0$ independent of the stated variables. Choose a cutoff $t_0 < e^{-e}$ within the range of the displayed bound, use its nonnegative error for $0 < t < t_0$, and set $r_c^{m,\kappa}(t) = 1$ for $t \geq t_0$. The latter range follows trivially from the product being at most one. Cauchy–Schwarz under any coupling with the specified marginals proves (A.9). $\square$

Write $\mathcal{S}_q = \{p \in [0, 1]^q : \sum_j p_j = 1\}$ and define the multilinear extension of $\overline{P}_0$ on $[q]^k$ by

$$P_0^{\text{ml}}(p_1, \dots, p_k) = \sum_{r \in [q]^k} \overline{P}_0(r) \prod_{i=1}^k p_i^{r_i}, \quad p_i \in \mathcal{S}_q. \quad (\text{A.10})$$

Lemma A.5 (A simplex-valued local estimate). *Let J, μ_c, α be the data of Lemma A.3, and fix $0 < \lambda < 1$. For an integer $R \geq 1$, choose c with probability ω_c . Conditional on this one choice of c , choose the R coordinate tuples independently from μ_c , producing strings $Z_1, \dots, Z_k \in [m]^R$. Then obtain $\tilde{Z}_i$ by independently resampling every position and coordinate uniformly with probability λ . For a single function $F = (F^1, \dots, F^q) : [m]^R \rightarrow \mathcal{S}_q$, used for every c and every position, suppose that $0 < t < 1/(kq)$ and*

$$\text{Inf}_\ell(T_{1-\lambda}F^j) \leq t \quad (\ell \in [R], j \in [q]).$$

Then

$$\mathbb{E}_{c, Z, \tilde{Z}} P_0^{\text{ml}}(F(\tilde{Z}_1), \dots, F(\tilde{Z}_k)) \leq \text{opt}_q(J) + r_s(kqt), \quad (\text{A.11})$$

where $r_s(u) \rightarrow 0$ as $u \downarrow 0$, independently of R . The function r_s depends only on the fixed local data, k, q, λ .

Proof. This is the simplex-valued estimate of [17, Lemma 10.5], with its sufficient scalar influence bounds and their dependence made explicit. Put $H = T_{1-\lambda}F$. Independence of the noise between positions and multilinearity turn the left side of (A.11) into

$$\mathbb{E}_c \mathbb{E}_{Z \sim \mu_c^{\otimes R}} P_0^{\text{ml}}(H(Z_1), \dots, H(Z_k)). \quad (\text{A.12})$$

For each coordinate, the integral ensemble consists of the indicators $\mathbf{1}_{\{X_i=j\}}$ under μ_c ; its atoms have probability greater than α . The common Gram data (A.7) supply one Gaussian ensemble for all variables of J with the same first and second moments on every constraint, as in [17, Lemma 10.4]. Explicitly, for a standard Gaussian vector g in the span of the Gram vectors, the variables

$$G_{u,j} = \langle I, b_{u,j} \rangle + \langle b_{u,j} - \langle I, b_{u,j} \rangle I, g \rangle$$

have those moments. Use independent copies of this ensemble in the R coordinates. Each H^j is evaluated on these ensembles through its multilinear product expansion; equivalently, express an orthonormal basis on $[m]$ as linear combinations of the symbol indicators and use the same polynomial for the integral and Gaussian ensembles.

For a fixed c , let $\mathcal{T}_\theta^c$ retain each whole coordinate tuple with probability θ and otherwise resample it from μ_c . Uniform one-position marginals give, componentwise,

$$\mathcal{T}_{1-\lambda}^c[F^j(Z_i)] = (T_{1-\lambda}F^j)(Z_i) = H^j(Z_i).$$

Thus these are the smoothed components required by the invariance theorem on the product of the tuple ensembles. The same marginals give the exact influence identity below; the superscript indicates the product measure used to compute influence:

$$\text{Inf}_\ell^{\mu_c^{\otimes R}}(H^j(Z_i)) = \text{Inf}_\ell^{[m]^R}(H^j). \quad (\text{A.13})$$

Indeed, conditioning on every coordinate tuple except the ℓ th leaves the ℓ th symbol at position i uniform, and the other symbols at that position have the uniform product law. For a vector, influence is the sum of its component influences [21, Definition 2.6]. Thus the vector $(H^j(Z_i))_{i \in [k], j \in [q]}$ has each coordinate influence at most kqt ; each component has variance at most one.

For the remaining tail hypothesis, let $(F^j(Z_i))_S$, $S \subseteq [R]$, be the Hoeffding components on the tuple product $\mu_c^{\otimes R}$, and let the superscript $> d$ retain the terms with $|S| > d$. The tuple-noise identity above multiplies component S by $(1 - \lambda)^{|S|}$. Since the unsmoothed F^j takes values in $[0, 1]$, for every real $d > 0$,

$$\begin{aligned} \text{Var}((H^j(Z_i))^{>d}) &= \sum_{|S|>d} (1 - \lambda)^{2|S|} \|(F^j(Z_i))_S\|_2^2 \\ &\leq (1 - \lambda)^{2d} \text{Var}(F^j(Z_i)) \leq (1 - \lambda)^{2d}. \end{aligned}$$

All norms and variances here use $\mu_c^{\otimes R}$. This supplies the tail bound in [21, Theorem 3.6]. These are sufficient hypotheses for the vector-valued invariance estimate stated in [17, Theorem 10.1].

Let $\mathbf{p} : \mathbb{R}^q \rightarrow \mathcal{S}_q$ be Euclidean projection. It is nonexpansive. The function

$$\Psi(x_1, \dots, x_k) = P_0^{\text{ml}}(\mathbf{p}(x_1), \dots, \mathbf{p}(x_k))$$

is Lipschitz with constant at most $\sqrt{kq}$: on the simplex, telescoping the positions bounds the payoff difference by $\sum_i \|x_i - y_i\|_1$, and projection and Cauchy–Schwarz give the claim. Apply the cited invariance estimate to this function, separately for each c , with vector dimension kq and influence threshold kqt . Its displayed error is at most

$$r_s(kqt), \quad r_s(u) = C_{kq} \sqrt{kq} u^{\lambda/(18 \log(1/\alpha))}.$$

In particular the error is independent of R .

On the integral ensemble, $H(Z_i) \in \mathcal{S}_q$, so Ψ equals the payoff in (A.12). On the common Gaussian ensemble, project $H(G_u)$ to $\mathcal{S}_q$ for every $u \in U$, and independently draw one q -label for

each u with those probabilities. This is one random assignment of J . The variables in every local tuple are distinct, so its conditional expected payoff at that tuple is the multilinear payoff Ψ . Every realized assignment has value at most $\text{opt}_q(J)$. Averaging the invariance comparisons over c proves (A.11). $\square$

Proof of Proposition A.2. We first derive estimates that hold for every R , then choose their numerical parameters independently of R . Fix for now $0 < \kappa < 1/2$, q with $1/q \leq \mu_0(\kappa)$, and local data from Lemma A.3. Use the local experiment of Lemma A.5 with the actual resampling probability $\lambda = 2\kappa$. Also fix auxiliary constants $0 < t < 1/(kq)$ and $\chi > 0$.

For an input graph Φ , let $\pi_{a \rightarrow b}$ denote the permutation on edge ab . Its action on strings is fixed by

$$(\pi_{a \rightarrow b} z)_{\pi_{a \rightarrow b}(\ell)} = z_\ell \quad (\ell \in [R]).$$

The output variable set will be $\mathcal{V} = B \times [m]^R$. Consider the following independent neighbor experiment: choose a uniformly from A , choose $b_1, \dots, b_k$ independently and uniformly from $N(a)$, and independently perform the local experiment, obtaining $\tilde{Z}_1, \dots, \tilde{Z}_k$. Its raw output positions are

$$w_i = (b_i, \pi_{a \rightarrow b_i} \tilde{Z}_i) \in \mathcal{V}. \quad (\text{A.14})$$

For any rank map $f : \mathcal{V} \rightarrow \mathbb{Z}$, define

$$V(f) = \mathbb{E} \overline{P}_0(f(w_1), \dots, f(w_k)), \quad V_a(f) = \mathbb{E}[\overline{P}_0(f(w_1), \dots, f(w_k)) \mid a].$$

The bounded payoff (A.4) is used even if the experiment repeats an output ID. We will obtain a lower bound for one coarse rank map and an upper bound for every injective order, then transfer both bounds to distinct-variable constraints.

Suppose a labeling L strongly satisfies at least $1 - \zeta$ of the left vertices. Define the coarse rank map $f(b, z) = z_{L(b)}$. At every strongly satisfied a and every neighbor b ,

$$f(b, \pi_{a \rightarrow b} z) = z_{\pi_{a \rightarrow b}^{-1}(L(b))} = z_{L(a)}.$$

Thus all k positions read the same label coordinate, for every choice of neighbors. With probability $(1 - 2\kappa)^k$ none of these k symbols is resampled; conditional on that event their tuple has law μ_c . Nonnegativity of the payoff gives

$$V(f) \geq (1 - \zeta)(1 - 2\kappa)^k v_{\text{loc}}. \quad (\text{A.15})$$

For soundness fix an injective order $O : \mathcal{V} \rightarrow [|\mathcal{V}|]$. At a fixed a , its local set $N(a) \times [m]^R$ has exactly $M = hm^R$ distinct IDs. Sort this set by O and divide the sorted list into q consecutive blocks $C_{a,1}, \dots, C_{a,q}$ of size M/q . These sizes are integers because $q \mid h$. Each block is the intersection of the local set with an interval of global ranks. Define

$$F_a^j(z) = \frac{1}{h} \sum_{b \in N(a)} \mathbf{1}_{\{(b, \pi_{a \rightarrow b} z) \in C_{a,j}\}}, \quad F_a = (F_a^1, \dots, F_a^q). \quad (\text{A.16})$$

For every z , $F_a(z) \in \mathcal{S}_q$. Each edge permutation is a bijection on strings, so

$$\mathbb{E}_z F_a^j(z) = \frac{|C_{a,j}|}{hm^R} = \frac{1}{q}. \quad (\text{A.17})$$

For any interval I of integer ranks put

$$F_a^I(z) = \frac{1}{h} \sum_{b \in N(a)} \mathbf{1}_{\{O(b, \pi_{a \rightarrow b} z) \in I\}}.$$

Call a *pseudorandom at threshold t* if $\text{Inf}_\ell(T_{1-\kappa}F_a^I) \leq t$ for every I and every $\ell \in [R]$. At such an a , this bound holds for each block function F_a^j .

Let $j_a(w)$ be the block index of a local ID w . Conditional on the local query strings, independence of the neighbor choices gives the exact identity

$$\mathbb{E}_{b_1, \dots, b_k} \overline{P}_0(j_a(w_1), \dots, j_a(w_k)) = P_0^{\text{ml}}(F_a(\tilde{Z}_1), \dots, F_a(\tilde{Z}_k)). \quad (\text{A.18})$$

The identity includes outcomes with equal neighbors. Further smoothing decreases the block influences, so $\text{Inf}_\ell(T_{1-2\kappa}F_a^j) \leq t$. Lemma A.5, applied to the single function F_a , bounds the expected coarse payoff in (A.18) by $\text{opt}_q(J) + r_s(kqt)$.

When the k coarse ranks are distinct, their induced permutation is the one induced by the fine ranks $O(w_i)$. Otherwise the two payoffs differ by at most one. This also covers a repeated output ID, which necessarily produces a coarse tie. For two positions $i \neq i'$, independence of their noise and neighbor choices gives their coarse tie probability as

$$\sum_{j=1}^q \mathbb{E}_c \mathbb{E}_{Z \sim \mu_c^{\otimes R}} [(T_{1-2\kappa}F_a^j)(Z_i)(T_{1-2\kappa}F_a^j)(Z_{i'})].$$

Conditional on c , both strings have uniform product marginals. Their correlation is permitted by Lemma A.4. Use that lemma with the fixed mean $1/q$ from (A.17), and write $r_c = r_c^{m, \kappa}$. The last display is at most $q^{-\kappa/2} + qr_c(t)$. A union bound over pairs therefore gives, at every pseudorandom a ,

$$V_a(O) \leq S_{\text{loc}} := \text{opt}_q(J) + \binom{k}{2} (q^{-\kappa/2} + qr_c(t)) + r_s(kqt). \quad (\text{A.19})$$

We next turn an excess over this bound into a labeling of Φ . For each $b \in B$, restrict O to $\{b\} \times [m]^R$ and compress its ranks to obtain a strict order o_b of $[m]^R$. Let

$$\mathcal{L}_b = \left\{ \ell \in [R] : \text{Inf}_\ell(T_{1-\kappa} \mathbf{1}_{\{o_b(z) \in I\}}) > t/2 \text{ for some rank interval } I \right\}.$$

The few-influential-coordinates lemma [17, Definition 4.2 and Lemma 4.3] gives

$$|\mathcal{L}_b| \leq \frac{400}{\kappa(t/2)^3} = \frac{3200}{\kappa t^3}. \quad (\text{A.20})$$

This bound is uniform in R .

If a is not pseudorandom, choose one witnessing interval I_a and coordinate ℓ_a with $\text{Inf}_{\ell_a}(T_{1-\kappa}F_a^{I_a}) > t$. Noise commutes with coordinate permutations and averaging, and influence is a convex squared norm. Hence

$$\text{Inf}_{\ell_a}(T_{1-\kappa}F_a^{I_a}) \leq \frac{1}{h} \sum_{b \in N(a)} \text{Inf}_{\pi_{a \rightarrow b}(\ell_a)} \left(T_{1-\kappa} \mathbf{1}_{\{O(b, z) \in I_a\}} \right). \quad (\text{A.21})$$

Each term is at most one. An average greater than t forces more than a $t/2$ fraction of the terms to exceed $t/2$. The restriction of a global-rank interval to one right code is an interval of its relative order, so these neighbors satisfy $\pi_{a \rightarrow b}(\ell_a) \in \mathcal{L}_b$.

Give each such a label ℓ_a , and label every other left vertex arbitrarily. Independently at each right vertex with $\mathcal{L}_b \neq \emptyset$, choose one uniform label from $\mathcal{L}_b$; use any label when the set is empty. By (A.20), the expected satisfied fraction among the edges at each nonpseudorandom a is greater than

$$\frac{t}{2} \frac{\kappa t^3}{3200} = \frac{\kappa t^4}{6400}.$$

One random right labeling is used for all left vertices. Linearity of expectation and the common left degree identify the resulting average with the ordinary uniform edge fraction.

If $V(O) > S_{\text{loc}} + \chi$, more than a χ fraction of left vertices are nonpseudorandom: if that fraction is g , the bound $0 \leq V_a(O) \leq 1$ and (A.19) give $V(O) \leq (1 - g)S_{\text{loc}} + g \leq S_{\text{loc}} + g$. The preceding decoder would then give a labeling of edge value greater than $\chi\kappa t^4/6400$. Consequently

$$\text{val}(\Phi) \leq \nu < \frac{\chi\kappa t^4}{6400} \implies V(O) \leq S_{\text{loc}} + \chi \quad \text{for every injective } O. \quad (\text{A.22})$$

It remains to turn the independent experiment into strict ordering constraints. Let D be the event that $b_1, \dots, b_k$ are distinct. For $h \geq k$,

$$p_D := \Pr(D) = \frac{(h)_k}{h^k}, \quad \delta_h := 1 - p_D \leq \frac{\binom{k}{2}}{h}. \quad (\text{A.23})$$

The probability is the same at every a , and D is independent of the local experiment. Conditional on D , a remains uniform and its ordered distinct neighbor tuple is uniform among its $(h)_k$ possibilities. All output IDs are then distinct because their first coordinates b_i differ. Let I_0 denote this finite weighted distribution of strict P_0 constraints, and write $\text{OPT}_{P_0}(I_0)$ for its maximum strict payoff.

Let $V_D(f)$ denote the expected tie-extended payoff after this conditioning, and use $V_{-D}(f)$ for conditioning on its complement. For any rank map f , coarse or injective,

$$V(f) = p_D V_D(f) + (1 - p_D) V_{-D}(f), \quad |V_D(f) - V(f)| \leq \delta_h, \quad (\text{A.24})$$

because all payoffs lie in $[0, 1]$. Apply this first to the coarse map in (A.15). Within each finite fiber of that map, choose one uniform random permutation of the distinct IDs, independently between fibers, and concatenate the fibers in rank order. On every conditioned tuple, restriction of these permutations to its IDs is uniform within each tied rank class. Thus the expected strict P_0 -payoff of this one random global order equals $V_D(f)$. Some global order attains at least that value. For soundness, an injective O has distinct ranks on every conditioned tuple, so its tie-extended payoff there is its strict P_0 -payoff. Equations (A.15), (A.22), and (A.24) therefore give the strict bounds

$$\begin{aligned} \text{YES : } \quad & \text{OPT}_{P_0}(I_0) \geq (1 - \zeta)(1 - 2\kappa)^k v_{\text{loc}} - \delta_h, \\ \text{NO : } \quad & \text{OPT}_{P_0}(I_0) \leq S_{\text{loc}} + \chi + \delta_h \quad \text{if } \nu < \chi\kappa t^4/6400. \end{aligned} \quad (\text{A.25})$$

These bounds are for the conditioned P_0 tuples. Applying (A.5) makes every occurrence use the original predicate Π with exactly the same values. Call this emitted instance I .

We now make the choices for the prescribed δ . Choose a positive rational η with $10\eta < \delta$, and a rational $0 < \kappa < 1/2$ with $2\kappa k < \eta$. Put $K = \binom{k}{2}$. Choose an integer q so large that

$$\frac{1}{q} \leq \mu_0(\kappa), \quad Kq^{-\kappa/2} < \eta.$$

Choose the data J, m, α from Lemma A.3 at error η . The two analytic error functions now have fixed parameters and are uniform in R . Choose a positive rational $t < 1/(kq)$ so small that

$$Kqr_c(t) < \eta, \quad r_s(kqt) < \eta.$$

Set $\chi = \eta$. Choose positive rationals $\zeta, \nu < 1$ with

$$\zeta < \eta, \quad \nu < \frac{\eta\kappa t^4}{6400},$$

and choose an integer $H \geq \max\{k, q\}$ with $K/H < \eta$. These choices depend only on k, Π, δ . In particular, they precede both the choice of h in Lemma A.1 and the alphabet supplied by Theorem 1.1.

For every $h \geq H$, $\delta_h < \eta$. The local gap (A.6), Bernoulli's inequality, and (A.25) give

$$\text{YES : } \text{OPT}_{\Pi}(I) \geq 1 - \zeta - 2\kappa k - \eta - \delta_h > 1 - 4\eta > 1 - \delta,$$

whereas (A.19) gives

$$\text{NO : } \text{OPT}_{\Pi}(I) \leq S_{\text{loc}} + \chi + \delta_h < \rho_{\Pi} + 6\eta < \rho_{\Pi} + \delta.$$

This proves the numerical claims simultaneously for every R .

Finally we specify the encoded transformation. Once R is fixed, the choice of c , the product $\mu_c^{\otimes R}$, and the independent rational 2κ -noise form one fixed finite rational law θ_R on local outcomes ξ . These are the actual local outcomes above; the Gaussian variables used for the analytic comparison are not part of the output. Renumber B by $[|B|]$ and explicitly list $\mathcal{V} = B \times [m]^R$. For each $a \in A$, each ordered distinct neighbor tuple, and each $\xi \in \text{supp } \theta_R$, list the tuple in (A.14), permuted by (A.5), with weight

$$\frac{\theta_R(\xi)}{|A|(h)_k}. \quad (\text{A.26})$$

The weights sum to one and are nonnegative rationals. Every listed tuple has distinct IDs; the same tuple may occur in several entries. There are $|B|m^R$ variables and at most $|A|(h)_k |\text{supp } \theta_R|$ entries. For fixed k, m, R, θ_R , these counts are polynomial in the explicit graph size, even when h varies. A fixed common denominator for θ_R shows that each weight has $O(\log |A| + k \log h)$ bits plus a fixed constant. The IDs have $O(\log |B|)$ bits plus a fixed constant. Exact enumeration of these entries is therefore deterministic and polynomial for fixed R , as asserted. $\square$

The same occurrence-list encoding also justifies the deterministic random-order guarantee in Section 8. Build an order by exposing one next variable at a time. If exactly s variables of a constraint have been exposed, its conditional success probability is the number of permutations in Π whose first s positions are the roles of those variables in the exposed order, divided by $(k-s)!$. It can be computed by inspecting the fixed set S_k . For each choice of the next variable, sum these probabilities against the listed rational weights and choose a next variable whose conditional expectation is at least the average over the choices. This preserves the initial expectation $|\Pi|/k!$. All comparisons are exact and polynomial in the input bit length: the product of the weight denominators has bit length at most their total encoded length, and the factorial denominators have constant size for fixed k . The variable list is explicit, so the number of prefix steps and candidate choices is polynomial as well.

References

- [1] Sanjeev Arora, Boaz Barak, and David Steurer. Subexponential algorithms for Unique Games and related problems. In *Proceedings of the 51st Annual IEEE Symposium on Foundations of Computer Science*, pages 563–572. IEEE Computer Society, 2010. doi:10.1109/FOCS.2010.59. <https://www.dsteurer.org/paper/subexpug.pdf>.
- [2] Sanjeev Arora, Carsten Lund, Rajeev Motwani, Madhu Sudan, and Mario Szegedy. Proof verification and the hardness of approximation problems. *Journal of the ACM*, 45(3):501–555, 1998. Preliminary version in FOCS 1992. doi:10.1145/278298.278306. <https://people.csail.mit.edu/madhu/papers/1992/almss-journ.pdf>.

- [3] Sanjeev Arora and Shmuel Safra. Probabilistic checking of proofs: A new characterization of NP. *Journal of the ACM*, 45(1):70–122, 1998. Preliminary version in FOCS 1992. doi:10.1145/273865.273901. <https://www.cs.umd.edu/~gasarch/TOPICS/pcp/AS.pdf>.
- [4] Mitali Bafna, Boaz Barak, Pravesh K. Kothari, Tselil Schramm, and David Steurer. Playing Unique Games on certified small-set expanders. In *Proceedings of the 53rd Annual ACM SIGACT Symposium on Theory of Computing*, pages 1629–1642. ACM, 2021. doi:10.1145/3406325.3451099. <https://arxiv.org/abs/2006.09969v3>.
- [5] Boaz Barak, Parikshit Gopalan, Johan Håstad, Raghu Meka, Prasad Raghavendra, and David Steurer. Making the long code shorter. *SIAM Journal on Computing*, 44(5):1287–1324, 2015. Preliminary version in FOCS 2012. doi:10.1137/130929394. <https://www.csc.kth.se/~johanh/shortcode.pdf>.
- [6] Boaz Barak, Pravesh K. Kothari, and David Steurer. Small-set expansion in shortcode graph and the 2-to-2 conjecture. In *10th Innovations in Theoretical Computer Science Conference (ITCS 2019)*, volume 124 of *Leibniz International Proceedings in Informatics*, pages 9:1–9:12. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2019. First circulated in 2018. Section-based theorem locators refer to arXiv:1804.08662v1. doi:10.4230/LIPIcs.ITCS.2019.9. <https://arxiv.org/abs/1804.08662v1>.
- [7] Mihir Bellare, Oded Goldreich, and Madhu Sudan. Free bits, PCPs, and nonapproximability—towards tight results. *SIAM Journal on Computing*, 27(3):804–915, 1998. doi:10.1137/S0097539796302531. <https://people.seas.harvard.edu/~madhusudan/papers/1995/bgs-journ.pdf>.
- [8] Moses Charikar, Konstantin Makarychev, and Yury Makarychev. Near-optimal algorithms for Unique Games. In *Proceedings of the Thirty-Eighth Annual ACM Symposium on Theory of Computing*, pages 205–214. ACM, 2006. doi:10.1145/1132516.1132547. <https://konstantin.makarychev.net/pdf/unique.pdf>.
- [9] Shuchi Chawla, Robert Krauthgamer, Ravi Kumar, Yuval Rabani, and D. Sivakumar. On the hardness of approximating multicut and sparsest-cut. *Computational Complexity*, 15(2):94–114, 2006. Numbered locators refer to the author manuscript. doi:10.1007/s00037-006-0210-9. <https://www.wisdom.weizmann.ac.il/~robi/papers/CKKRS-CutHardness-CompComp06.pdf>.
- [10] Erik D. Demaine, Dotan Emanuel, Amos Fiat, and Nicole Immorlica. Correlation clustering in general weighted graphs. *Theoretical Computer Science*, 361(2–3):172–187, 2006. Numbered locators refer to the author manuscript dated February 5, 2005. doi:10.1016/j.tcs.2006.05.008. https://erikdemaine.org/papers/Clustering_TCS/paper.pdf.
- [11] Irit Dinur, Subhash Khot, Guy Kindler, Dor Minzer, and Muli Safra. On non-optimally expanding sets in Grassmann graphs. *Israel Journal of Mathematics*, 243(1):377–420, 2021. First circulated in 2017; preliminary version in STOC 2018. doi:10.1007/s11856-021-2164-7.
- [12] Irit Dinur, Subhash Khot, Guy Kindler, Dor Minzer, and Muli Safra. Towards a proof of the 2-to-1 games conjecture? *Theory of Computing*, 21(11):1–50, 2025. Preliminary version in STOC 2018. doi:10.4086/toc.2025.v021a011. <https://theoryofcomputing.org/articles/v021a011/v021a011.pdf>.

- [13] Irit Dinur and David Steurer. Analytical approach to parallel repetition. In *Proceedings of the Forty-Sixth Annual ACM Symposium on Theory of Computing*, pages 624–633. ACM, 2014. Numbered locators refer to the author manuscript dated June 10, 2014. doi:10.1145/2591796.2591884. <https://www.dsteurer.org/paper/productgames.pdf>.
- [14] David Ellis, Guy Kindler, and Noam Lifshitz. An analogue of Bonami’s lemma for functions on spaces of linear maps, and 2-2 games, 2026. First circulated in 2022; version 2, January 22, 2026. <https://arxiv.org/abs/2209.04243v2>.
- [15] Uriel Feige, Shafi Goldwasser, László Lovász, Shmuel Safra, and Mario Szegedy. Interactive proofs and the hardness of approximating cliques. *Journal of the ACM*, 43(2):268–292, 1996. doi:10.1145/226643.226652. <https://www.cs.tau.ac.il/~safra/PapersAndTalks/FGLSS.pdf>.
- [16] Michel X. Goemans and David P. Williamson. Improved approximation algorithms for maximum cut and satisfiability problems using semidefinite programming. *Journal of the ACM*, 42(6):1115–1145, 1995. doi:10.1145/227683.227684.
- [17] Venkatesan Guruswami, Johan Håstad, Rajsekar Manokaran, Prasad Raghavendra, and Moses Charikar. Beating the random ordering is hard: Every ordering CSP is approximation resistant. *SIAM Journal on Computing*, 40(3):878–914, 2011. Numbered locators refer to ECCC TR11-027. doi:10.1137/090756144. <https://ecc.ecc.ecc.ecc.ac.il/report/2011/027/download>.
- [18] Johan Håstad. Some optimal inapproximability results. *Journal of the ACM*, 48(4):798–859, 2001. Theorem locators refer to the author manuscript dated February 18, 2002. doi:10.1145/502090.502098. <https://people.kth.se/~johanh/optimalinap.pdf>.
- [19] Wassily Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963. doi:10.1080/01621459.1963.10500830.
- [20] Thomas Holenstein. Parallel repetition: Simplifications and the no-signaling case. *Theory of Computing*, 5(8):141–172, 2009. doi:10.4086/toc.2009.v005a008. <https://theoryofcomputing.org/articles/v005a008/v005a008.pdf>.
- [21] Marcus Isaksson and Elchanan Mossel. Maximally stable gaussian partitions with discrete applications. arXiv:0903.3362v3 [math.PR], 2009. Numbered locators refer to arXiv:0903.3362v3 (3 August 2009). <https://arxiv.org/pdf/0903.3362v3>.
- [22] Subhash Khot. On the power of unique 2-prover 1-round games. In *Proceedings of the Thirty-Fourth Annual ACM Symposium on Theory of Computing*, pages 767–775. ACM, 2002. doi:10.1145/509907.510017.
- [23] Subhash Khot. On the Unique Games Conjecture. In *Proceedings of the 25th Annual IEEE Conference on Computational Complexity*, pages 99–121. IEEE Computer Society, 2010. Conjecture locators refer to the author manuscript. doi:10.1109/CCC.2010.19. <https://cs.nyu.edu/~khot/papers/UGCSurvey.pdf>.
- [24] Subhash Khot, Guy Kindler, Elchanan Mossel, and Ryan O’Donnell. Optimal inapproximability results for MAX-CUT and other 2-variable CSPs? *SIAM Journal on Computing*, 37(1):319–357, 2007. Numbered theorem references use the author manuscript dated September 19, 2005. Author manuscript. doi:10.1137/S0097539705447372.

- [25] Subhash Khot, Dor Minzer, and Muli Safra. Pseudorandom sets in Grassmann graph have near-perfect expansion. *Annals of Mathematics*, 198(1):1–92, 2023. Theorem locators refer to ECCC TR18-006, revision 2, May 19, 2018; the expansion theorem is Theorem 1.12. doi:10.4007/annals.2023.198.1.1. <https://eccc.weizmann.ac.il/report/2018/006/revision/2/download>.
- [26] Subhash Khot, Dor Minzer, and Muli Safra. On independent sets, 2-to-2 games and Grassmann graphs. *Theory of Computing*, 21(10):1–55, 2025. Preliminary version in STOC 2017. doi:10.4086/toc.2025.v021a010.
- [27] Subhash Khot and Oded Regev. Vertex cover might be hard to approximate to within $2 - \epsilon$. *Journal of Computer and System Sciences*, 74(3):335–349, 2008. Numbered locators refer to the linked author manuscript. doi:10.1016/j.jcss.2007.06.019. https://cims.nyu.edu/~regev/papers/vc_hard.pdf.
- [28] Subhash Khot and Muli Safra. A two-prover one-round game with strong soundness. *Theory of Computing*, 9(28):863–887, 2013. Preliminary version in FOCS 2011. doi:10.4086/toc.2013.v009a028.
- [29] Subhash Khot and Nisheeth K. Vishnoi. The Unique Games Conjecture, integrality gap for cut problems and embeddability of negative-type metrics into ℓ_1 . *Journal of the ACM*, 62(1):8:1–8:39, 2015. Preliminary version in FOCS 2005. doi:10.1145/2629614. <https://arxiv.org/abs/1305.4581>.
- [30] Alexandra Kolla. Spectral algorithms for Unique Games. *Computational Complexity*, 20(2):177–206, 2011. Preliminary version in CCC 2010. doi:10.1007/s00037-011-0011-7. <https://arxiv.org/abs/1102.2300>.
- [31] Dor Minzer and Kai Zhe Zheng. Near optimal alphabet-soundness tradeoff PCPs, 2026. Version 4, revised May 13, 2026. Preliminary version in STOC 2024, pages 15–23, doi:10.1145/3618260.3649606. <https://arxiv.org/abs/2404.07441v4>.
- [32] Elchanan Mossel, Ryan O’Donnell, and Krzysztof Oleszkiewicz. Noise stability of functions with low influences: Invariance and optimality. *Annals of Mathematics*, 171(1):295–341, 2010. doi:10.4007/annals.2010.171.295.
- [33] OpenAI. Constant-factor hardness of uniform sparsest cut. OpenAI Math Release preprint OAI:Constant-factor-hardness-of-uniform-sparsest-cut-September-24-2026, 2026.
- [34] OpenAI. Perfect completeness for 2-to-1 games. OpenAI Math Release preprint OAI:Perfect-completeness-for-2-to-1-games-September-23-2026, 2026.
- [35] OpenAI. The Gaussian propeller bound in every dimension. OpenAI Math Release preprint OAI:The-Gaussian-Propeller-Bound-in-Every-Dimension-September-24-2026, 2026.
- [36] Prasad Raghavendra. Optimal algorithms and inapproximability results for every CSP? In *Proceedings of the Fortieth Annual ACM Symposium on Theory of Computing*, pages 245–254. ACM, 2008. doi:10.1145/1374376.1374414.
- [37] Prasad Raghavendra. *Approximating NP-hard Problems: Efficient Algorithms and Their Limits*. PhD thesis, University of Washington, 2009. <https://people.eecs.berkeley.edu/~prasad/Files/thesis.pdf>.

- [38] Anup Rao. Parallel repetition in projection games and a concentration bound. *SIAM Journal on Computing*, 40(6):1871–1891, 2011. Theorem 4 refers to the author manuscript dated September 17, 2010. doi:10.1137/080734042. <https://homes.cs.washington.edu/~anuprao/pubs/parrep.pdf>.
- [39] Ran Raz. A parallel repetition theorem. *SIAM Journal on Computing*, 27(3):763–803, 1998. Preliminary version in STOC 1995. doi:10.1137/S0097539795280895.
- [40] Jacob T. Schwartz. Fast probabilistic algorithms for verification of polynomial identities. *Journal of the ACM*, 27(4):701–717, 1980. doi:10.1145/322217.322225.
- [41] Luca Trevisan. Approximation algorithms for Unique Games. *Theory of Computing*, 4(5):111–128, 2008. Preliminary version in FOCS 2005. doi:10.4086/toc.2008.v004a005. <https://theoryofcomputing.org/articles/v004a005/v004a005.pdf>.
- [42] Richard Zippel. Probabilistic algorithms for sparse polynomials. In *Symbolic and Algebraic Computation (EUROSAM 1979)*, volume 72 of *Lecture Notes in Computer Science*, pages 216–226. Springer, 1979. doi:10.1007/3-540-09519-5_73.